

LESSON 10

Coding with Matlab

A complete example about the similarity in images



Lesson outline

- Concept of Cross-Correlation
- Matlab first steps with images
- Coding an example for similarity
- **Main points**







THEORY

Similarity in images via cross-correlation

What is relevant what is useless?






Similarity by Cross-correlation

- Is a measure of similarity of two vectors (matrices) as a function of the displacement of one relative to the other.
- The function returns
 - the **similarity** value
 - and the position with **the better alignment** of the two images.

`[similarity, displacement] = image_xcorr(I1, I2)`



Similarity by Cross-correlation

- The “core of cross-correlation”:
for image-processing applications in which the brightness of the image and template can vary due to lighting and exposure conditions, the images can be first normalized

- This is typically done at every step by subtracting the mean and dividing by the standard deviation.

Template = $t(x,y)$, Subimage $f(x,y)$

Repeat
shifting the
subimage f to get
the alignment points
checking where is
the maximum value
of NCC

Normalized
Cross-Correlation = $t * f = \frac{1}{n} \sum_{x,y} \frac{1}{\sigma_f \sigma_t} (f(x,y) - \mu_f) (t(x,y) - \mu_t).$
(NCC)


$f(x+i,y+j)$



Toolboxes

Matlab

First step with images and similarity



Knowing what to learn from the coding lesson?

- Matlab **first steps with images**
- Understanding the structure and steps of the coding example for the image similarity
- During the exam will be **no exact coding**
- It necessary to understand the main procedures (+ what are inputs and outputs)
- For example, today we will better understand the **cross-corr.**



Why similarity is **so important**?

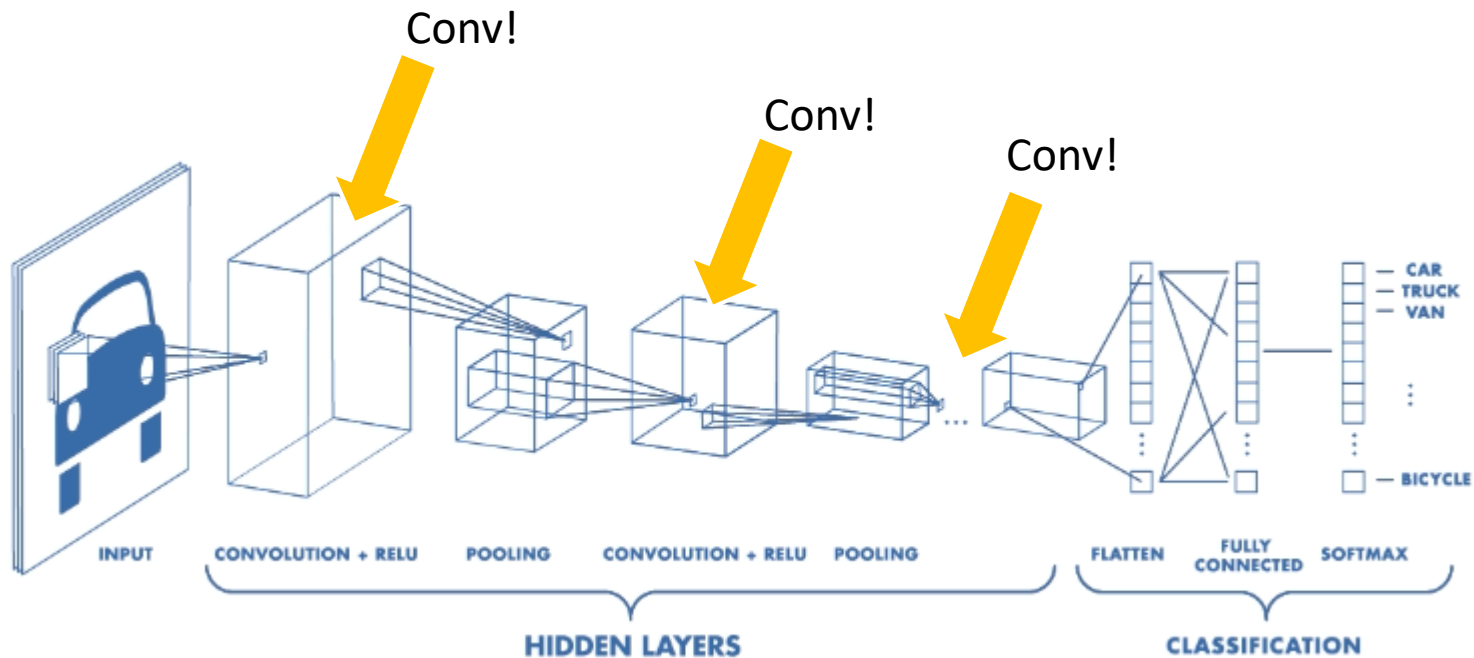


- Because it is one of the basis/kernel/primitive of the **pattern recognition** and **machine learning**
- **You/The Intelligent system can recognize a pattern/object/vector/image because it is similar to a previously labelled one seen the learning phase**

Why similarity is so important?




Convolution \leftrightarrow Correlation Convolutional Neural Networks



Have you noticed the puzzle?

(Lesson 9)


Image Similarity



- **Image Similarity** compares two images and returns a value that tells you how visually similar they are.
`score = similarity(image1, image2)`
- If Score $\rightarrow 0$ so images \rightarrow Contextually similar
- If (score == 0) so images are identical

Fabio Scotti - Università degli Studi di Milano

Did you find all the differences?
Are you sure?




How to automate the task?

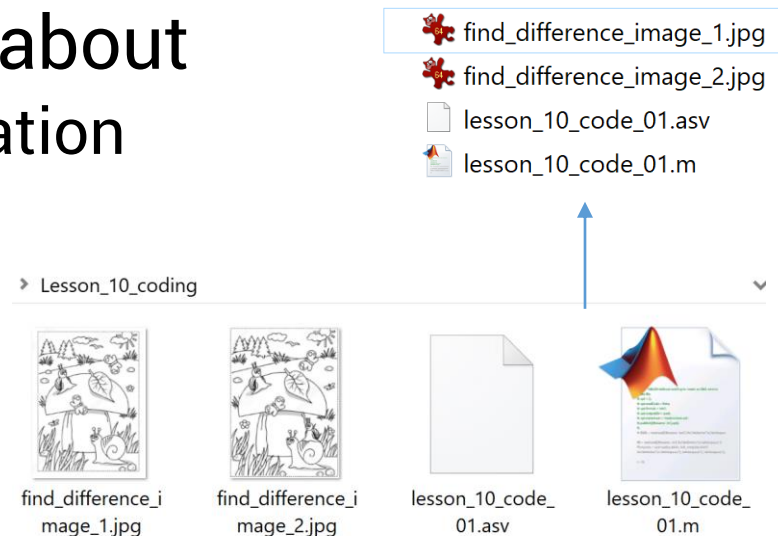
All what we need is

- Concept of image cross-correlation
- ...and a little of matlab coding

The material you need



- Matlab (last version)
- Download the files of the lesson in local folder
- Do not jump to the solution....
- Follow the process,
you will gather information about
 - Image format and representation
 - Image manipulation
 - Image plotting
 - Image similarity





Loading the files

Not compiled!

```
close all
clear all % do not use it in case of debug...

img_template = imread('find_difference_image_1.jpg');
img_subimage = imread('find_difference_image_2.jpg');
```



Name	Size	Bytes	Class
img_subimage	691x505x3	1046865	uint8
img_template	736x555x3	1225440	uint8

Color image
(RGB)

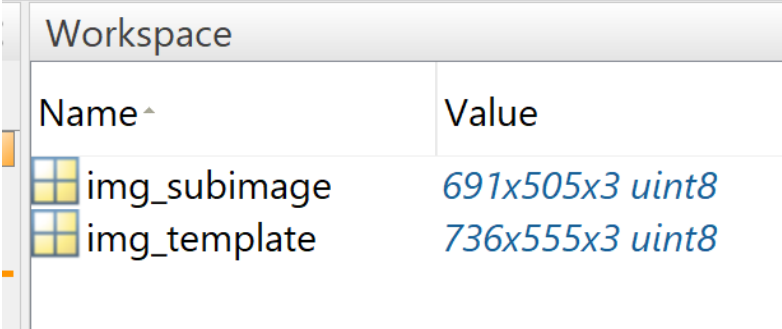
Unsigned
char

Unsigned 8-bit integers.



The uint8 range is from 0 to 255

Have a look to the variables

- The workspace
- `>> whos`



The image shows a screenshot of the MATLAB Workspace window. It has a title bar 'Workspace' and a table with two columns: 'Name' and 'Value'. There are two rows of variables: 'img_subimage' with value '691x505x3 uint8' and 'img_template' with value '736x555x3 uint8'. Each row has a small icon to its left.

Workspace	
Name ^	Value
 img_subimage	691x505x3 uint8
 img_template	736x555x3 uint8

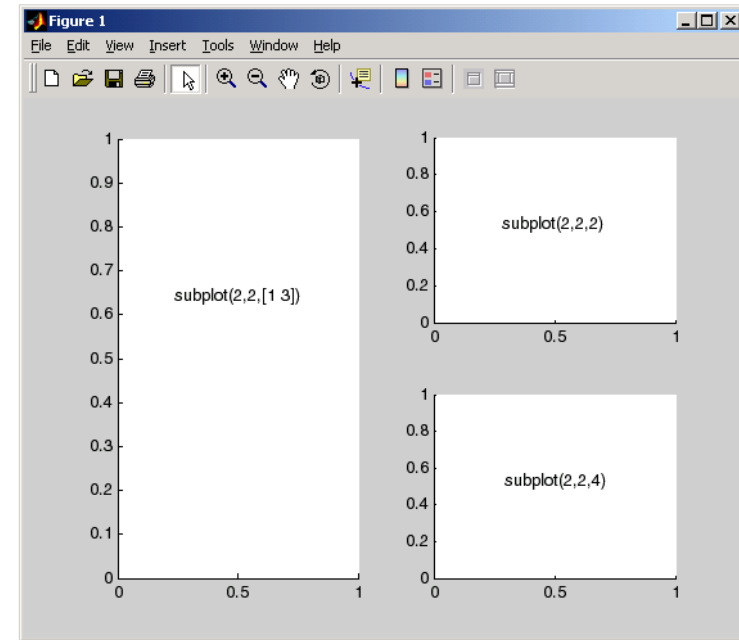
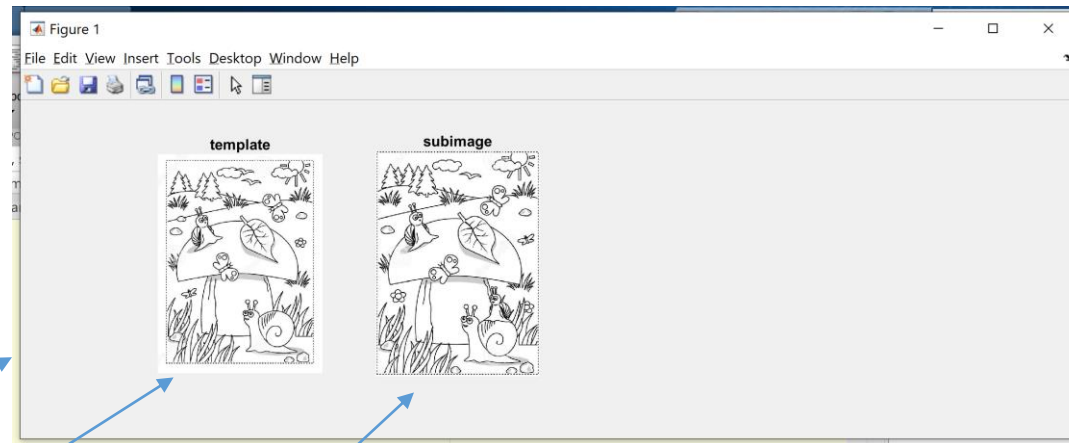
Name	Size	Bytes	Class
img_subimage	691x505x3	1046865	uint8
img_template	736x555x3	1225440	uint8

Plotting images

handle

```
h1 = figure;  
subplot(1,4,1)  
imshow(img_template, [])  
title('template')
```

```
subplot(1,4,2)  
imshow(img_subimage, [])  
title('subimage')
```



Plotting images (2)



```
h1 = figure;  
subplot(1,4,1)  
imshow(img_template, [])  
title('template')
```


```
subplot(1,4,2)  
imshow(img_subimage, [])  
title('subimage')
```

Auto color scale (**Square brackets**)

imshow(I,[]) displays the grayscale image I scaling the display based on the range of pixel values in I. **imshow** uses [min(I(:)) max(I(:))] as the display range, that is, the minimum value in I is displayed as black, and the maximum value is displayed as white.

How to see the difference?

The general idea is to subtract two images to find the differences



Name	Size	Bytes	Class
img_subimage	691x505x3	1046865	uint8
img_template	736x555x3	1225440	uint8

- 1) Images are different in size → we can't just subtract them (impossible in every coding language)
- 2) We have to place properly the second image to avoid false detections

Color conversion

```
% color 2 gray conversion  
img_template_gray = rgb2gray( img_template );  
img_subimage_gray = rgb2gray( img_subimage );
```

Name	Size	Bytes	Class
h1	1x1	8	matlab.ui.Figure
img_subimage	691x505x3	1046865	uint8
img_subimage_gray	691x505	348955	uint8
img_template	736x555x3	1225440	uint8
img_template_gray	736x555	408480	uint8

The handle of the figure points to h1.

gray images points to img_subimage_gray and img_template_gray.

Image normalization (mean subtraction)

```
img_template_gray_norm = img_template_gray - mean(mean(img_template_gray)) ;  
img_subimage_gray_norm = img_subimage_gray - mean(mean(img_subimage_gray)) ;
```

Just without the means (no standard deviations)

Two times..

mean Average or mean value.

$S = \text{mean}(X)$ is the mean value of the elements in X if X is a vector.
For matrices, S is a row vector containing the mean value of each column.

For N-D arrays, S is the mean value of the elements along the first array dimension whose size does not equal 1.


```
>> mean(mean(img_template_gray))
```

```
ans =
```

```
228.9717
```

It's quite white!

white == 255
@uint8



2D correlation (similarity evaluation)

GENERAL

Normalized cross-correlation = (NCC)

$$t * f = \frac{1}{n} \sum_{x,y} \frac{1}{\sigma_f \sigma_t} (f(x,y) - \mu_f) (t(x,y) - \mu_t).$$

MATLAB → with no scaling (direct!)

```
% xcorr2(A,B) computes the crosscorrelation of matrices A and B.
crr = xcorr2(img_template_gray_norm , img_subimage_gray_norm);
```

2-D Cross-Correlation

The 2-D cross-correlation of an M -by- N matrix, X , and a P -by- Q matrix, H , is a matrix, C , of size $M+P-1$ by $N+Q-1$.

$$C(k,l) = \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} X(m,n) \overline{H(m-k,n-l)}, \quad \begin{aligned} -(P-1) &\leq k \leq M-1, \\ -(Q-1) &\leq l \leq N-1, \end{aligned}$$

2D correlation (similarity evaluation)

```
% xcorr2(A,B) computes the crosscorrelation of matrices A and B.
crr = xcorr2(img_template_gray_norm , img_subimage_gray_norm);
```

Normalized cross-correlation = (NCC)

$$t * f = \frac{1}{n} \sum_{x,y} \frac{1}{\sigma_f \sigma_t} (f(x,y) - \mu_f) (t(x,y) - \mu_t).$$


Repeat shifting the subimage f to get the alignment points checking where is the maximum value of NCC $f(x+i,y+j)$



```
>> whos crr
```

Name	Size	Bytes	Class	Attributes
crr	1426x1059	12081072	double	

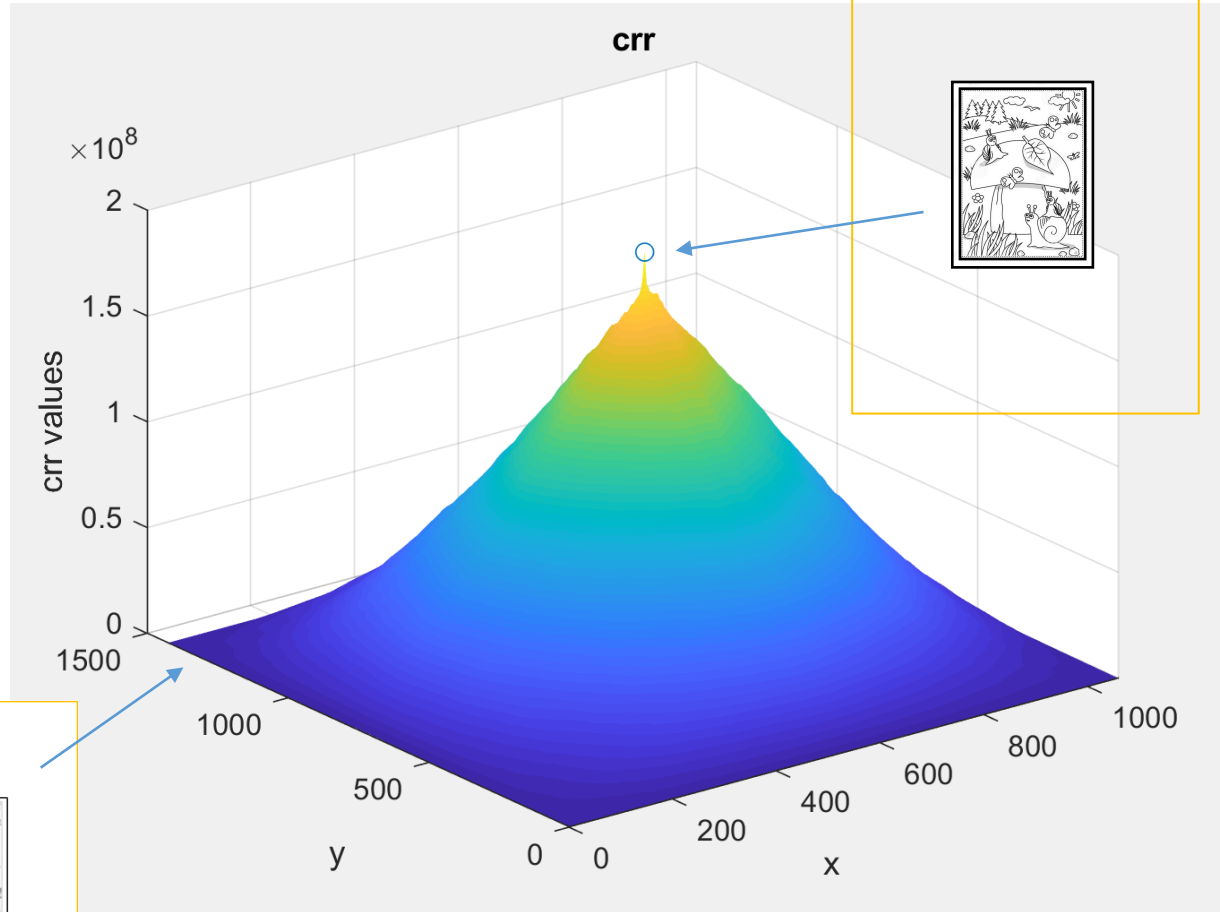
Name	Size	Bytes	Class
h1	1x1	8	matlab.ui.Figure
img_subimage	691x505x3	1046865	uint8
img_subimage_gray	691x505	348955	uint8
img_template	736x555x3	1225440	uint8
img_template_gray	736x555	408480	uint8



How is the output of a xcorr2?

```
%%  
figure  
surf(crr);  
shading interp;  
title('crr')
```

just to remove black
line around the tiles


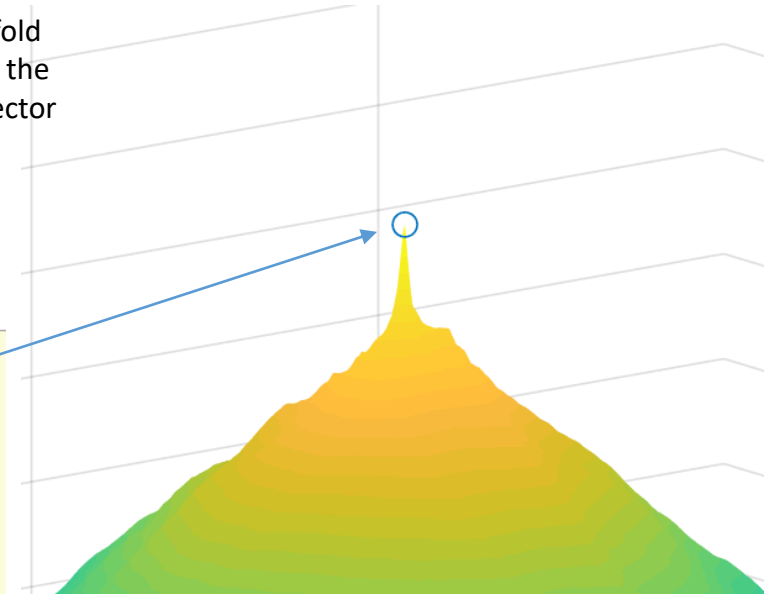


A real spike! Locate the max in ccr!

```
% find the maximum
[ssr,snd] = max( crr(:) );
[ij,ji] = ind2sub(size(crr),snd);
```

operator (:) unfold
the elements of the
matrix in a 1D vector

```
%%
hold on
plot3( ji, ij, ssr, 'o')
xlabel('x');
ylabel('y');
zlabel('crr values');
```



See the difference?

→ Image subtraction (Part 1)

size

The basic idea is to see the differences by subtracting the original template with a proper image containing the subimage properly **sized** and **shifted**

$$\text{img_delta} = \text{template} - \text{img_diff}$$

```
% copy the original image (**just for the EXTERNAL FRAME**)
img_diff = img_template_gray ;
```



Template

This is just to

- 1) Get the proper **dimensions** since `img_diff` must be subtracted
- 2) Copy the borders since the subimage is missing the borders to have a more readable comparison

See the difference?

→ Image subtraction (Part 2)

The proper shift

```
% Place the smaller image inside the larger image. Rotate the smaller image
% to comply with the convention that MATLAB® uses to display images.

img_diff(ij:-1:ij-size(img_subimage_gray,1)+1,ji:-1:ji-size(img_subimage_gray,2)+1) = ...
    rot90(img_subimage_gray,2) ;
```

```
% find the maximum
[ssr,snd] = max( crr(:) );
[ij,ji] = ind2sub(size(crr),snd);
```

rot90 Rotate array 90 degrees.


$B = \text{rot90}(A)$ is the 90 degree counterclockwise rotation of matrix A . If A is an N-D array, $\text{rot90}(A)$ rotates in the plane formed by the first and second dimensions.

$\text{rot90}(A,K)$ is the $K \times 90$ degree rotation of A , $K = +1, +2, \dots$

Example,

$A = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$	$B = \text{rot90}(A) = \begin{bmatrix} 3 & 6 \\ 2 & 5 \\ 1 & 4 \end{bmatrix}$
--	---

This notation is not in the exam!



Memory and
visualizations are
ordered differently


See the difference?

→ Image subtraction (check)


Open the previous figure by using the handle

```
figure(h1)
subplot(1,4,3)
imshow( img_diff )
title('just shift before subtraction')
```

template



subimage



just shift before subtraction




See the difference?

→ Image subtraction (check 1)

Open the previous figure by using the handle

```
figure(h1)
subplot(1,4,3)
imshow( img_diff )
title('just shift before subtraction')
```

OK We have the subimage inside **img_diff**



See the difference?


→ Image subtraction (check 2)

Open the previous figure by using the handle

```
figure(h1)
subplot(1,4,3)
imshow( img_diff )
title('just shift before subtraction')
```

OK We have the subimage inside

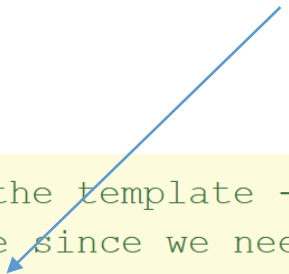
OK With the correct shifts



Creating the delta image

We need a double image not uint8!!

We need to be ready to store negative values in the delta!




```
% img_diff2 = the template - the shifted subimage (with just the frame of template)
% we use double since we need negative values (not allowed in the uint8 format)
img_delta = double(zeros(size(img_template_gray)));
img_delta = double(img_template_gray) - double(img_diff);
```

Creating the delta image (2)

We need to be ready to store negative values in the delta!

```
% img_diff2 = the template - the shifted subimage (with just the frame of template)
% we use double since we need negative values (not allowed in the uint8 format)
img_delta = double(zeros(size(img_template_gray)));
img_delta = double(img_template_gray) - double(img_diff);
```

Make a «cast»!
To avoid any
conversion
problems!



Checking the delta image

```
% figure(h1)
subplot(1,4,4)
imshow( img_delta, [] )
title('shift and subtraction')
```

template



subimage



just shift before subtraction



shift and subtraction



Checking the delta image




difference image



Improving the visualization

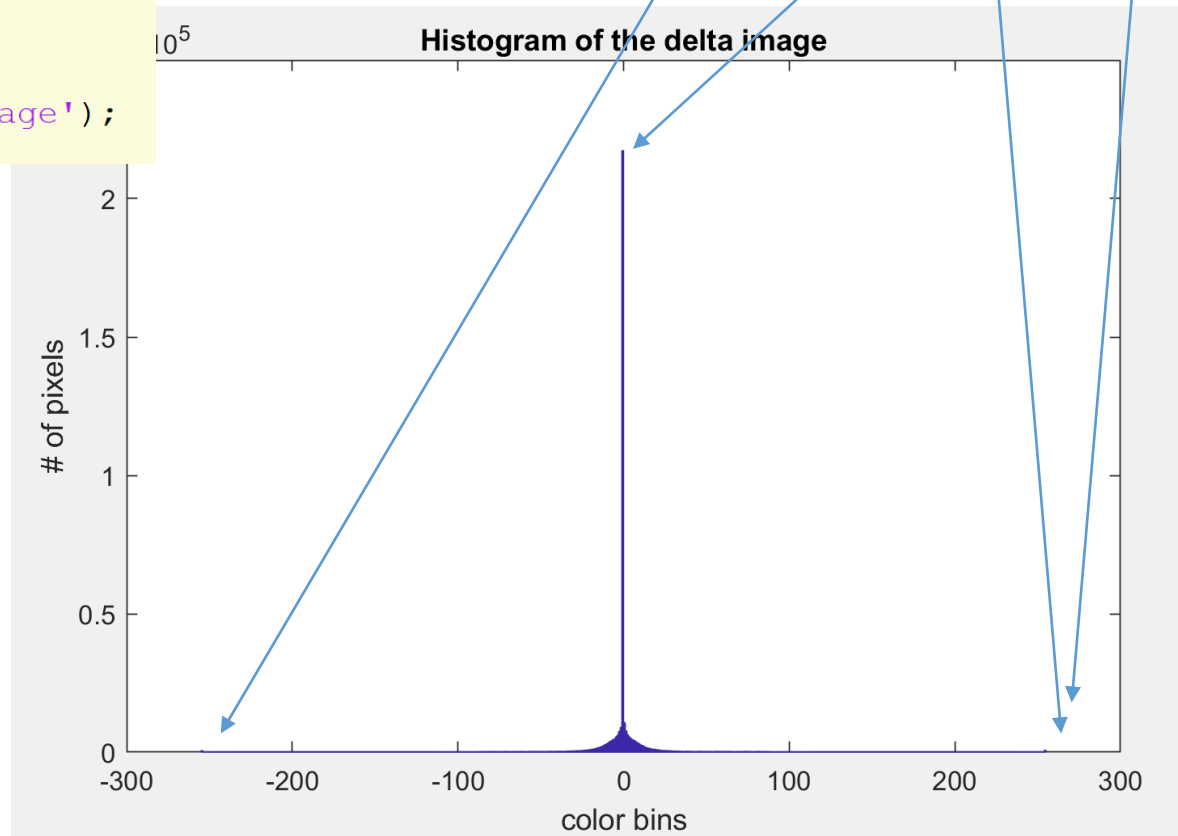

```
% improving the visualization
figure
subplot(1,2,1)
hist( img_delta(:), [2*255]);
xlabel('color bins')
ylabel('# of pixels')
title('Histogram of the delta image');
```

Understanding
how to use the
image histogram
in your work!




Sometimes histograms
are too “flat”
to really understand
what is happening

Zero values are grey
(histogram stretching)



Improving the visualization

```
subplot(1,2,2)
[h_y, h_x]= hist( img_delta(:), [2*255]);
subplot(1,2,2)
semilogy(h_x, h_y);
xlabel('color bins')
ylabel('# of pixels')
title('Logaritimic histogram of the delta image');
```



Separating the additions and deletions


```
%% finding the good thresholds
```

```
max_value = max(img_delta(:));
```

```
min_value = min(img_delta(:));
```

```
additions = (img_delta > 0.5 * max_value) ;
```

```
deletions = (img_delta < 0.5 * min_value) ;
```



Separating the additions and deletions

```
%% finding the good thresholds
```

```
max_value = max(img_delta(:));
```


```
min_value = min(img_delta(:));
```

```
additions = (img_delta > 0.5 * max_value) ;
```

```
deletions = (img_delta < 0.5 * min_value) ;
```

```
>> whos additions
```

Name	Size	Bytes	Class
additions	736x555	408480	logical



Separating the additions and deletions


```
max_value = max(img_delta(:));
min_value = min(img_delta(:));

additions = (img_delta > 0.5 * max_value) ;
deletions = (img_delta < 0.5 * min_value) ;


figure
subplot(1,3,1)
imshow( img_delta , [])
title('difference image')

subplot(1,3,2)
imshow( additions , [])
title('additions')


subplot(1,3,3)
imshow( deletions , [])
title('deletions')
```




difference image




additions




deletions




Overview




additions



deletions



Homeworks



- How is changing the CRR if you put in input the **same template image** as subimage
 - What about the maximum?
 - What about the shape of the CRR?
- How is changing the CRR if you put **two completely different images** in input?
 - What about the maximum?
 - What about the shape of the CRR?

Main points



- Cross-Correlation
- Matlab first steps with images
- Coding an example for similarity