

Intelligent systems for industry, supply chain and environment

LESSON 16

In case of new models.

Linear regressors and classifiers (with coding)



Outline

- *New model?*

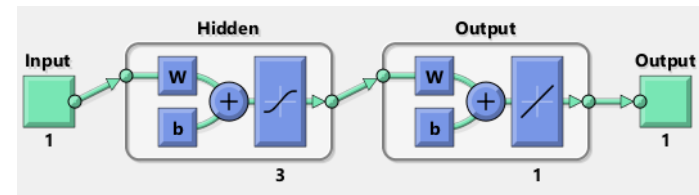
- *What to know*
- *What to do*



- *Creation and use (with code) of the first complete machine learning models*

- *Linear regressors*
- *Linear classifiers*
- *First neural networks*

- *Main points*

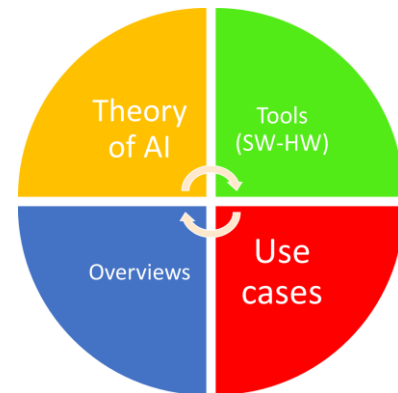




THEORY

New model?

Things to know for a correct start



What to know in case of new models?

- *Everyday, new models/learning techniques will be available. How to face this challenge and keep the pace?*
- **Consult documentation and examples** before to waste time
 - **Understand the topic, the type, the context!**
 - Check on what libraries models be used
 - General or specific?
 - What the computational complexity of the model?
 - In train
 - In test/deployment
 - Need CUDA?



«Let's use it, we don't need to read documentation!»

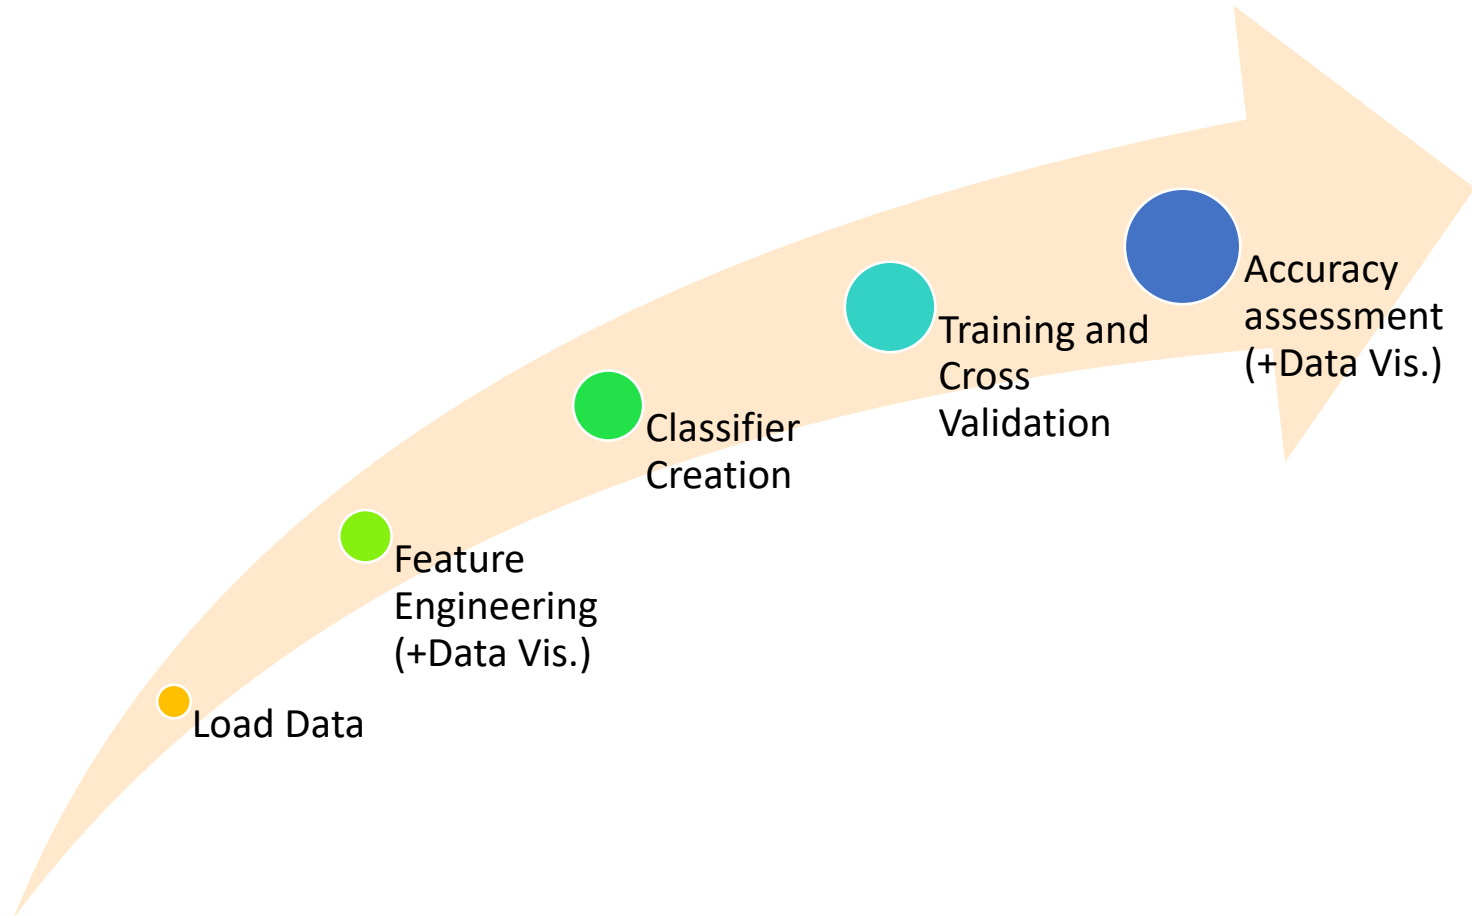
What to know in case of new models?

- Consult documentation and examples before to waste time (cont...)
 - **Memory** needs
 - Training
 - Test
 - Examples/**Dataset size**
1000 sample or 10M samples?
 - **Parameters**
 - Complex parameters to understand?
 - Sensitive to parameter tuning?
 - Can you do just a **fine-tuning** or need to re-train?
 - **Portability** on different libraries/environments
 - Is it open/GNU/free or just one single company are running it?

What to do in case of new models?

- *Try demos/examples in the documentation before to waste time, checking if it is/not suitable for your applications*
- *Try new models on a set of standard datasets and compare the accuracy and computational times*
 - *Not just the IRIS dataset..*
- *Remember to Occam's razor!*
 - *Same performance, but more complex, does it worth it?*

Application of a new classifier

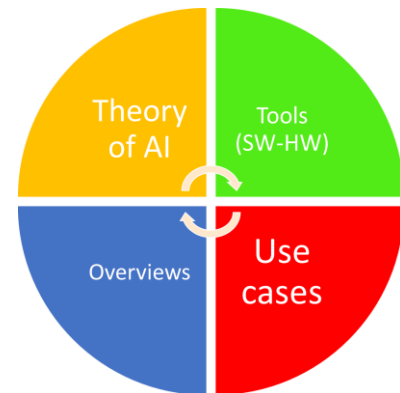




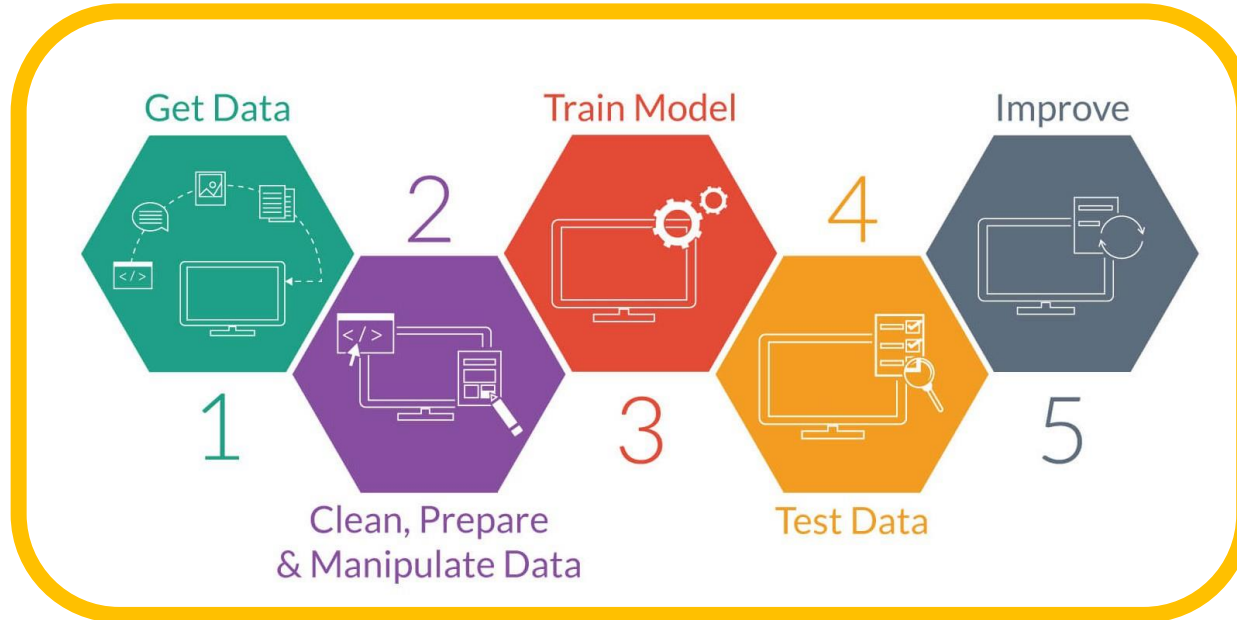
THEORY

A simple example of the complete workflow

The linear regressor

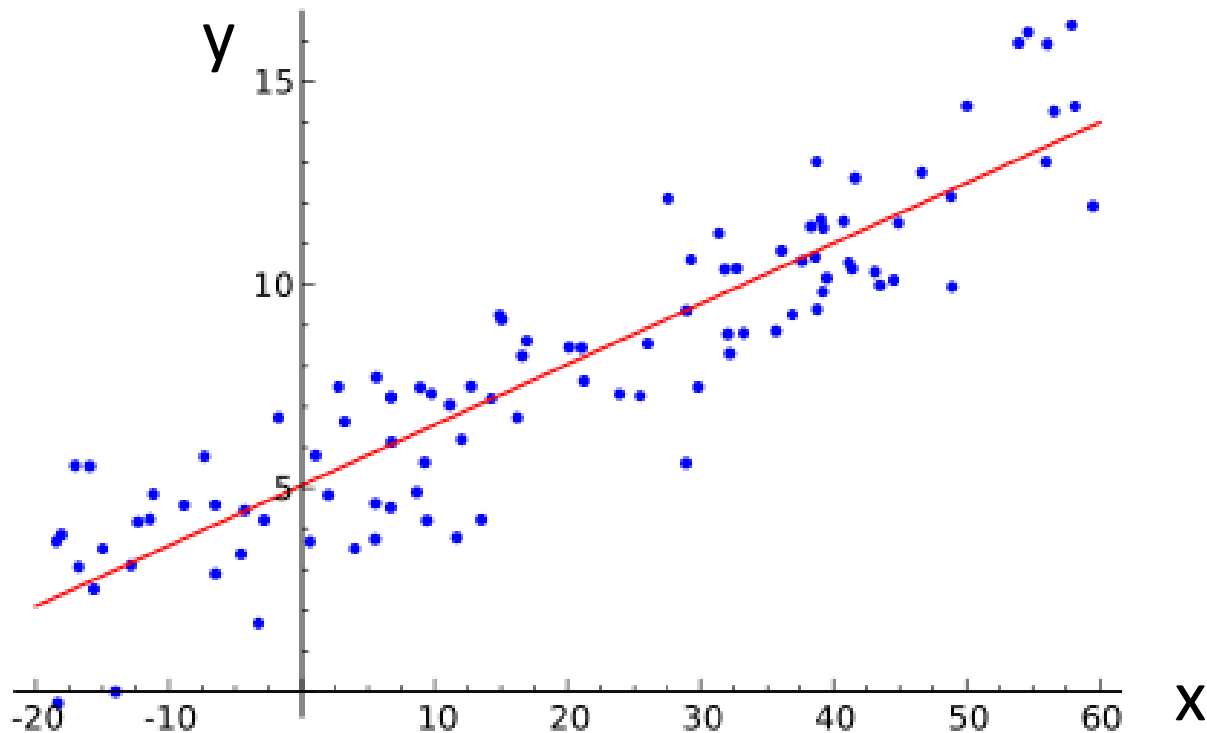


Let's review all the steps of the desing workflow with a simple example



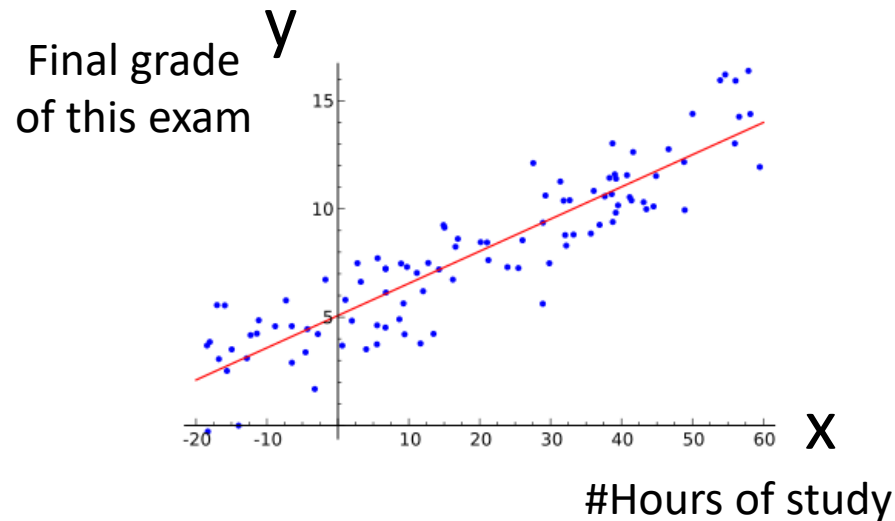
Linear Regression

We want to find the best linear function $y=f(x)$ to explain the data we have



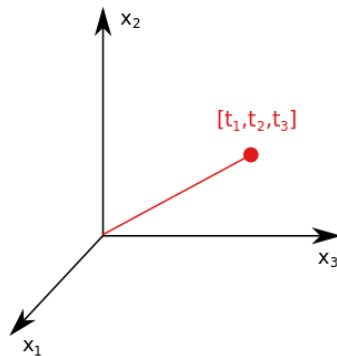
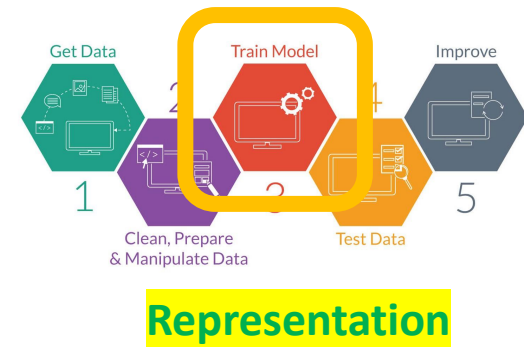
Linear Regression examples

- Age, gender, and diet \rightarrow height of the kid
- Type and grades of college \rightarrow yearly income
- One more..



An interesting example...

Extendible to N features



Vectorial form

$$z = \mathbf{w} \cdot \mathbf{x} + b$$

$$z = [w_1 \ w_2 \ w_3] \cdot \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} + b$$

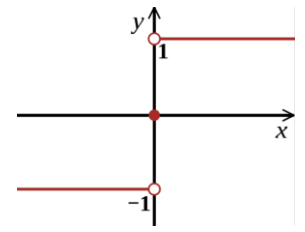
To completely describe the model
you need to fix 4 parameters:

w_1, w_2, w_3, b .

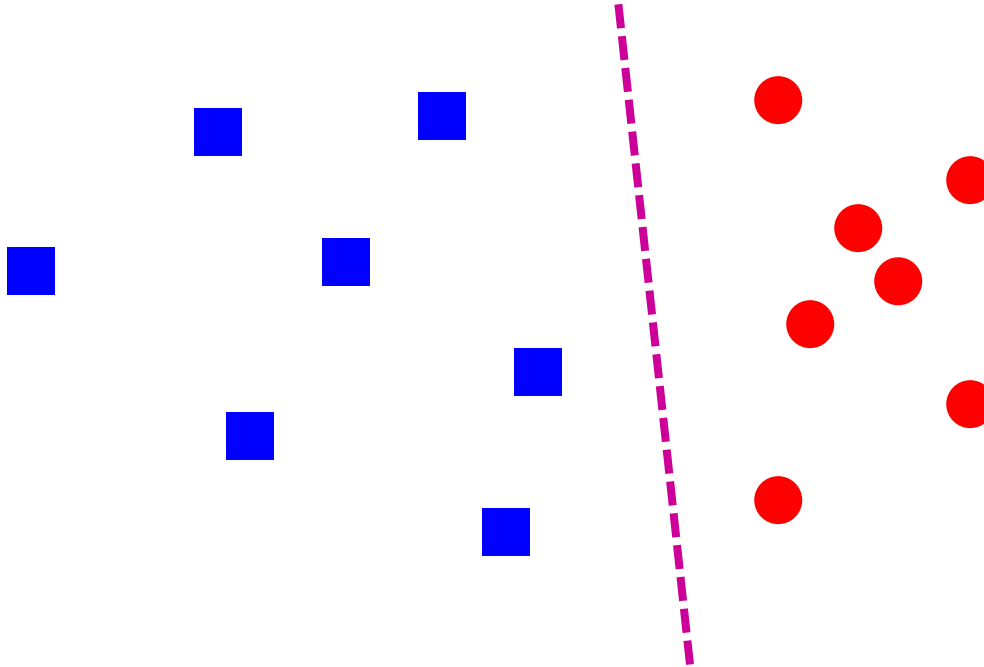
→ You need 4 data points

But you will use much more data!

From regressor to classifiers: Linear Class.



$$\text{sgn}(x) := \begin{cases} -1 & \text{if } x < 0, \\ 0 & \text{if } x = 0, \\ 1 & \text{if } x > 0. \end{cases}$$



Problem

Find a linear function to separate the classes:

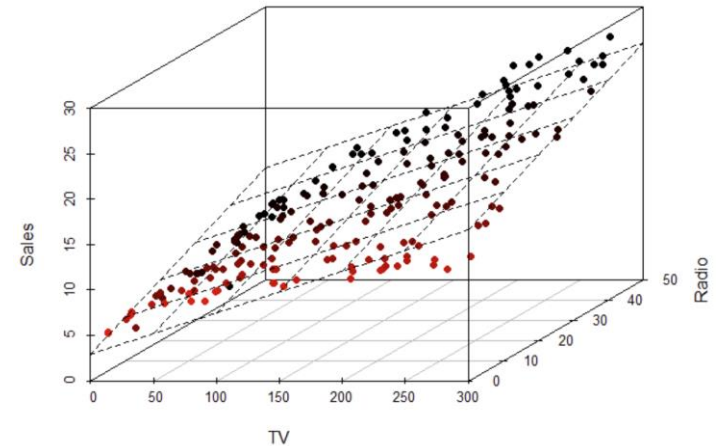
$$f(x) = \text{sgn}(\underbrace{w \cdot x + b}_{\text{Linear regressor}})$$

Linear regressor

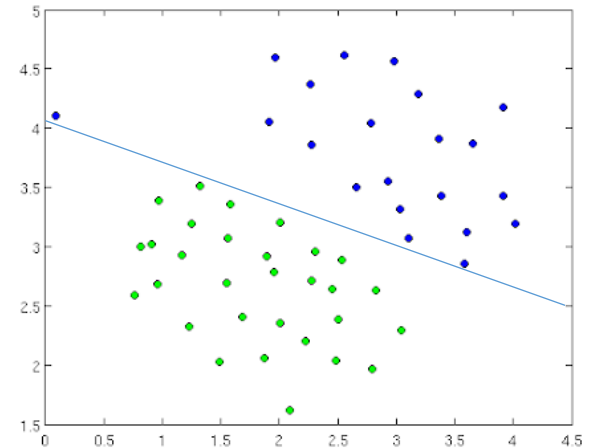
Linear regressor
 $f_1(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x} + b$

Linear classifier:
 $f_2(\mathbf{x}) = \text{sgn}(\underbrace{\mathbf{w} \cdot \mathbf{x} + b}_{\text{Linear regressor}})$

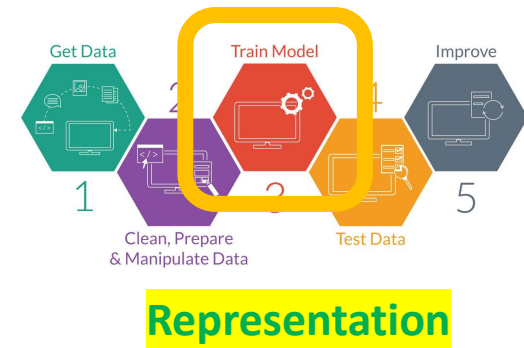
2D example



2D example



Linear Regression



The predicted value of y is given by:

$$\hat{y} = \hat{\beta}_0 + \sum_{j=1}^p X_j \hat{\beta}_j$$

The vector of coefficients $\hat{\beta}$ is the regression model.

JUST A NOTE about computation:

If you add a term in the inputs $X_0 = 1$, the formula becomes a simple matrix product (a compact notation): $\hat{\mathbf{y}} = \mathbf{X} \hat{\boldsymbol{\beta}}$

The algebra to get the solution is more easy to deal with.

About notation...

Elements

$$\hat{y} = \hat{\beta}_0 + \sum_{j=1}^p X_j \hat{\beta}_j$$

Vectorial form

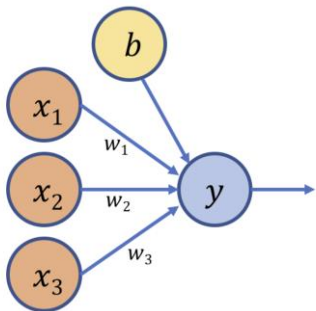
$$\mathbf{z} = \mathbf{w} \cdot \mathbf{x} + \mathbf{b}$$

«**weights**»

Since they multiply
the information coming
from the inputs

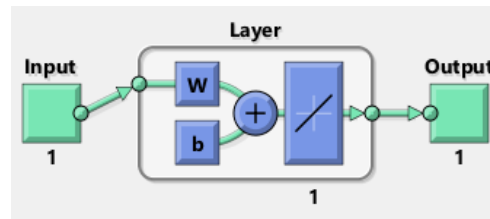
«**bias**»

Something
always present,
not modulated
from the input



That is (also)
the linear output
of a neuron

3 inputs



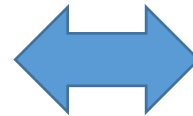
1 inputs

Linear Regression, #Features>1

$$f(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x} + b$$

We can write all of the input samples from the training in a single matrix \mathbf{X} :

$$\mathbf{X} = \begin{pmatrix} x_{11} & \cdots & x_{1n} \\ \vdots & \ddots & \vdots \\ x_{m1} & \cdots & x_{mn} \end{pmatrix}$$



	Individual #1								Individual #n
Age,	5.1								
gender,	3.5								
diet,	1.4								
....	0.2								

n = # samples (**distinct observations**)

m = # features

$$\mathbf{y} = (y_1, y_2, \dots, y_n)$$



Glucose

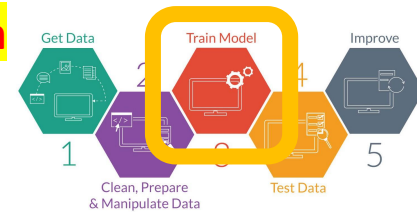
3.0								
-----	--	--	--	--	--	--	--	--

Vector \mathbf{Y} does not contain
classes here!!

We are creating a regressor
→ we will have numbers in \mathbf{Y}

Residual Sum-of-Squares

Optimization



To determine the model parameters $\hat{\beta}$ from some data, we minimize the Residual Sum of Squares which is

$$\text{RSS}(\beta) = \sum_{i=1}^N (y_i - \beta x_i)^2$$

Appendix (non in the exam) using the «compact notation» :

or symbolically $\text{RSS}(\beta) = (\mathbf{y} - \mathbf{X}\beta)^T (\mathbf{y} - \mathbf{X}\beta)$. To minimize it, take the derivative w.r.t. β which gives:

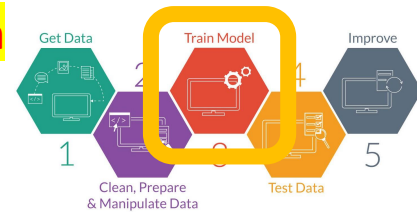
$$\mathbf{X}^T (\mathbf{y} - \mathbf{X}\beta) = 0$$

And if $\mathbf{X}^T \mathbf{X}$ is non-singular, the unique solution is:

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Iterative Regression Solutions

Optimization



Appendix (not in the exam):

- The exact method requires us to invert a matrix $(\mathbf{X}^T \mathbf{X})^{-1}$. This will often be **too big**.
- There are many **gradient-based** methods which reduce the RSS error by taking the **derivative wrt** β

$$\text{RSS}(\beta) = \sum_{i=1}^N (y_i - \beta x_i)^2$$

which was

$$\nabla = \mathbf{X}^T (\mathbf{y} - \mathbf{X}\beta)$$

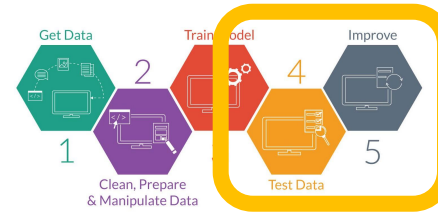
Important note!

Many neural method (even in deep learning) are using gradient-based methods.

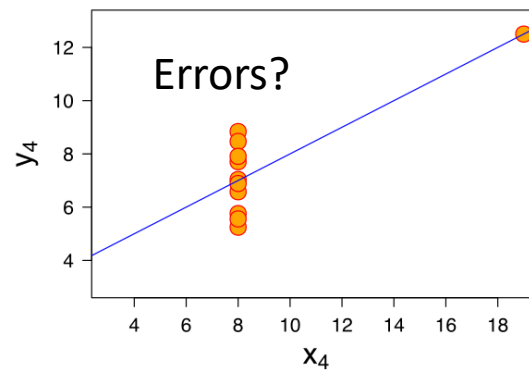
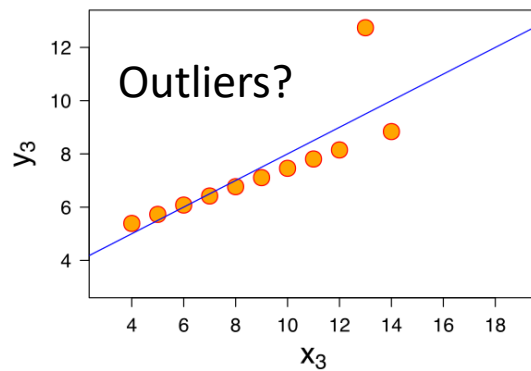
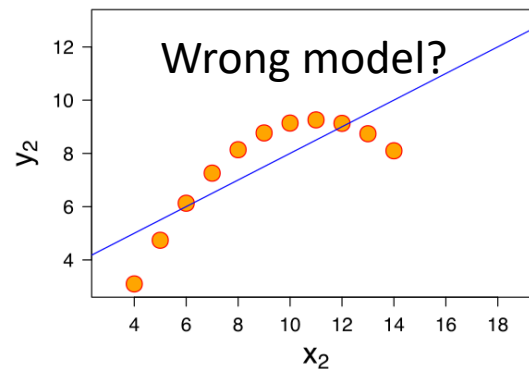
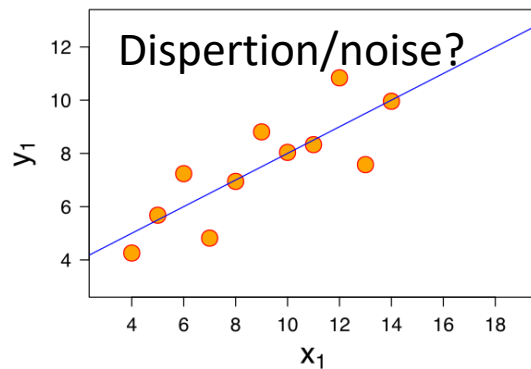
If the gradient $\rightarrow 0$ the method “gets stuck”

Is the regression good?

R^2 -value



- We can **always** fit a linear model to any dataset, but how do we know if there is a **real linear relationship**?



$R^2 = 0.67$ for all
Dim > 2 it is hard to
visually understand
the fitting quality!

Property	Value	Accuracy
Mean of x	9	exact
Sample variance of x : σ^2	11	exact
Mean of y	7.50	to 2 decimal places
Sample variance of y : σ^2	4.125	± 0.003
Correlation between x and y	0.816	to 3 decimal places
Linear regression line	$y = 3.00 + 0.500x$	to 2 and 3 decimal places, respectively
Coefficient of determination of the linear regression: R^2	0.67	to 2 decimal places

Anscombe's quartet: four data sets that have nearly identical simple descriptive statistics

R²-values (R-squared)

Let \hat{y} be a predicted value, and \bar{y} be the sample mean. Then the R-squared statistic is

$$R^2 = 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2}$$

And can be described as the fraction of the total variance not explained by the model.

R-squared

→ 1 means good fit

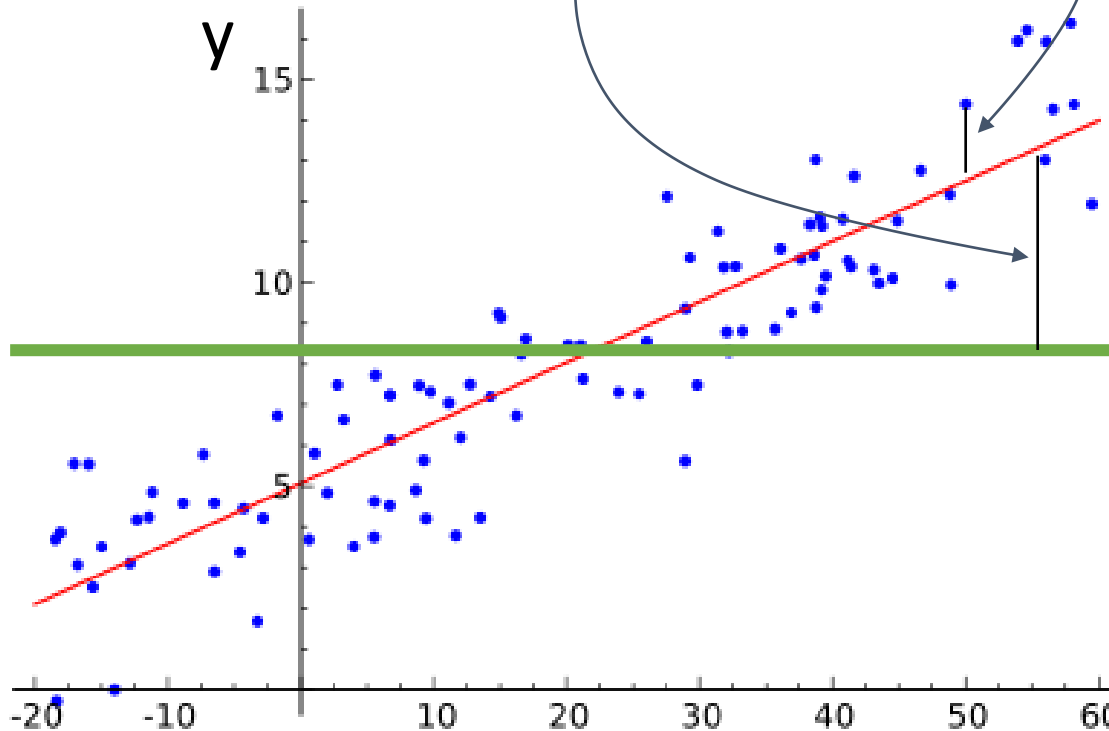
$$R^2 = 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2} =$$

sum of squares
of residuals

total sum of squares
(proportional to the
variance of the data)

Line of \hat{y}

Line of \bar{y} (mean of data)



R^2 It must be reported!

When Dim > 2

it is quite important

(you can't see...)

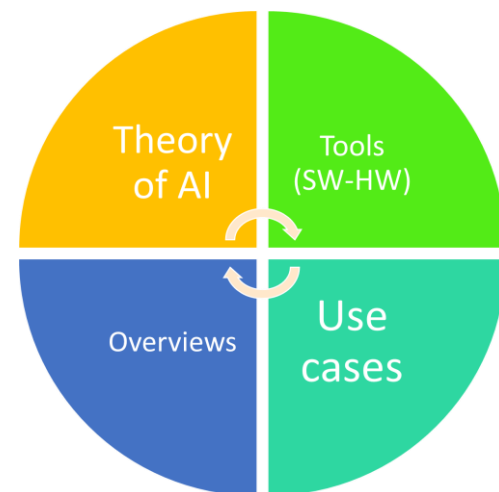


Toolboxes

Matlab and Python

Linear regressor

Linear classifiers



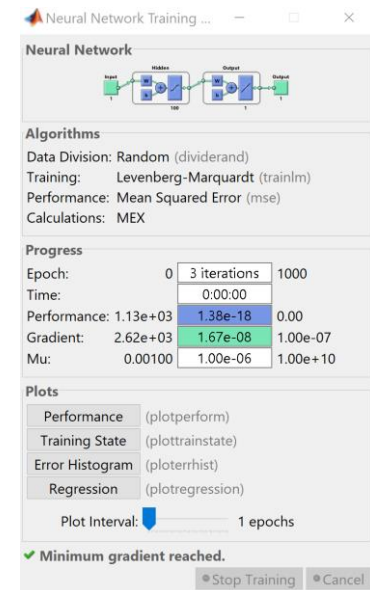
Matlab file

- 1) *Follow next slides*
- 2) *and then lunch your scrip.*



laboratory_MATLAB_linear_and_neural.m

Starting from regression
we will train and test
the first neural networks

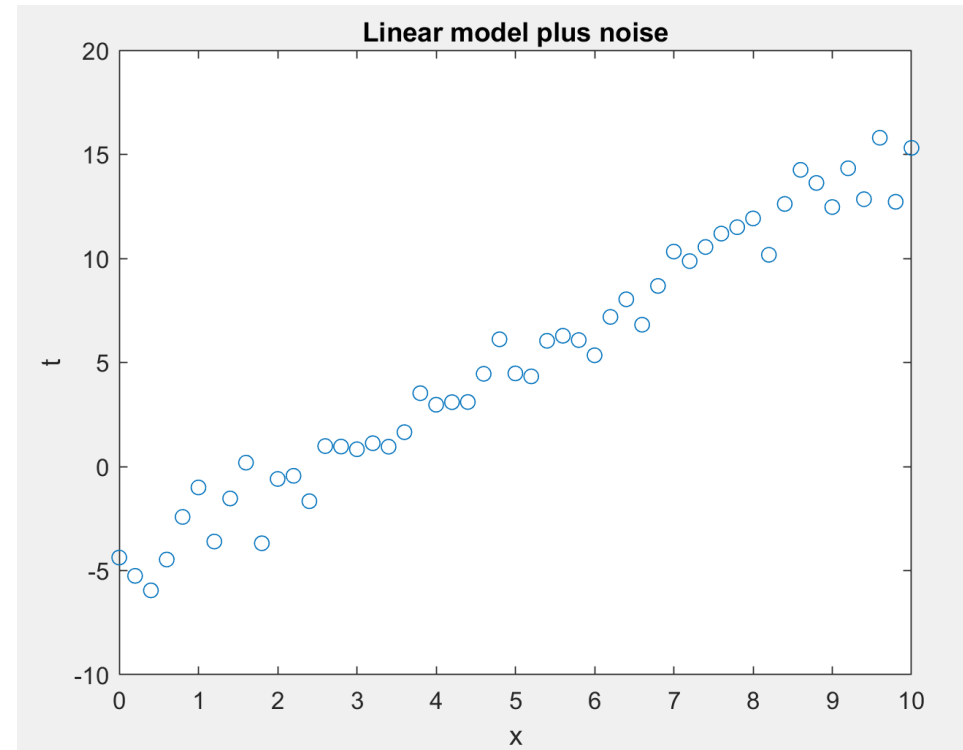


Matlab: create and train a linear regressor

Create 51 points for the eq. $t = 2x - 5$ plus gaussian noise

```
% creating a data with known distribution
% PROBLEM 01
x = [0:0.2:10];
noise = randn(size(x));
t = 2 * x - 5 + noise;

h1 = figure ;
plot(x,t, 'o');
title('PROBLEM 01 liner data plus noise');
xlabel('x');
ylabel('t');
```



Create and train a linear regressor

net1 = newlind(x,t)

METHOD#1 (direct, see Lesson #15 from data matrix)

```
% create the linear model and train it  
  
if (1) % optimized solution based on the inversion of the matrix of data  
    net1 = newlind(x,t);  
else % not optimized this is the general learning for non-linear networks
```

```
end
```

data partition

Create and train a linear regressor

You can create the linear regressor using the initialization and learning for neural networks

```
%% create the linear model and train it

if (1) % optimized solution based on the inversion of the matrix of data
    % net1 = perceptron;
else % not optimized this is the general learning for non-linear networks
    % net1 = perceptron;
    net1 = linearlayer;
    net1 = configure(net1,[0],0);

    % A little help to tune the inzial points
    net1.IW{1,1} = 3.1;
    net1.b{1} = -3.1

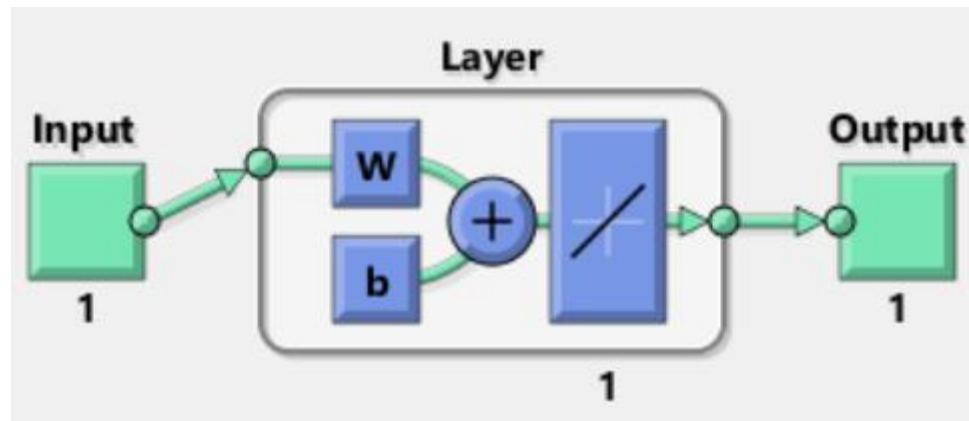
    [net1 , tr] = train(net1,x,t); % WARNING --> Train/Validation auto.created
    figure(h2); plotperform(tr) % need more tuning..... :-(
end
```

METHOD#2

- inzialization as linear net
- 1 input 1 layer
- Weight inzialization
- Levemberg-Marquardt optimization
- (general, but not optimized)

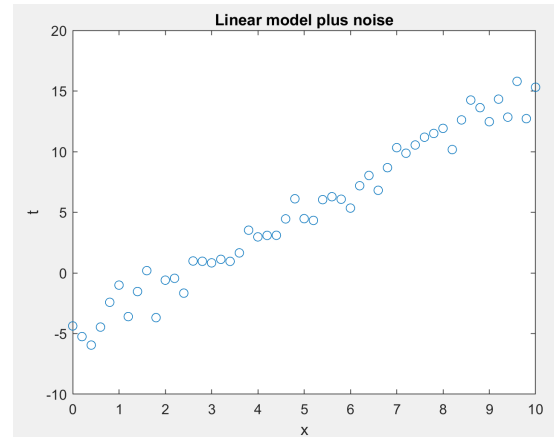
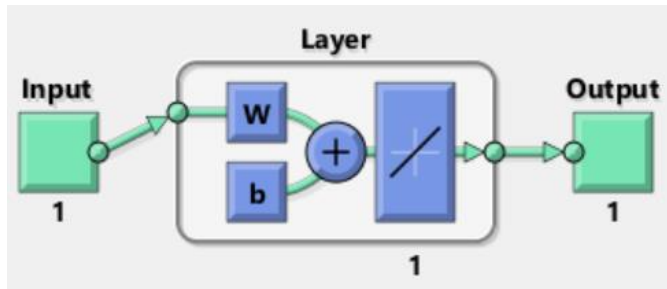
The created linear regressor

```
%% visualization of the trained linear model  
view(net1)
```



$$\text{Output} = w \text{ input} + b$$

Access the parameters



```
x = [0:0.2:10];  
noise = randn(size(x));  
t = 2 * x - 5 + noise;
```

```
% how to access the weight of the model  
W = net1.IW{1,1} ;  
b = net1.b{1} ;  
fprintf('-----\n');  
fprintf('trained linear model W = %f \n' , W);  
fprintf('trained linear model b = %f \n' , b);  
fprintf('-----\n');
```

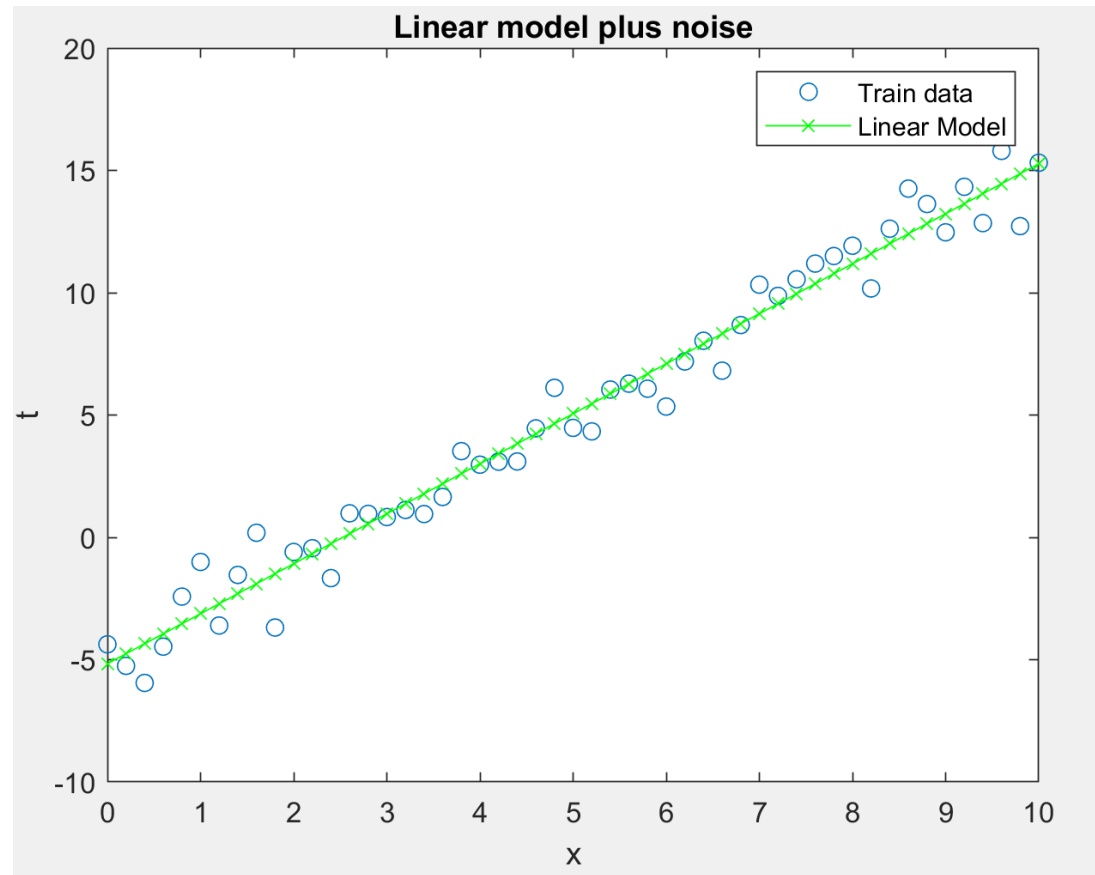
```
-----  
trained linear model W = 2.043733  
trained linear model b = -5.165622  
-----
```

Processing inputs and plotting the results

```
h1 = figure ;  
plot(x,t, 'o');  
title('Linear model plus noise');  
xlabel('x');  
ylabel('t');  
  
% how to check the performance  
y1 = net1(x); % simulate the inputs  
perf = perform(net1,y1,t) % Error on **training**  
  
% plot the output of the model  
figure(h1)  
hold on  
plot(x, y1, 'xg-')  
legend('Train data', 'Linear Model')
```

net1 = newlind(x,t)

y1 = net1(x)



Accuracy assessment

%% how to check the performance

```
y1 = net1(x); % simulate the inputs
perf = perform(net1,y1,t) % Error on **trainig**
net1.performFcn % is telling you the error metrics (MSE for Lin.)
mse_check = sum((y1 - t).^2 )/ size(t,2);
mean_t = mean(t);
R2 = 1 - ( sum((t - y1).^2 )/sum((t - mean_t).^2) );
fprintf('-----trained linear model-----\n');
fprintf('Perf.           = %f \n' , perf );
fprintf('Type of index = %s \n' , net1.performFcn );
fprintf('MSE              = %f \n' , mse_check );
fprintf('R2               = %f \n' , R2 );
fprintf('-----\n');
```

$$MSE = \sum_{i=1}^n \frac{(\hat{y}_i - y_i)^2}{n}$$

$$R^2 = 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2}$$

****WARNING****

This is about TRAIN dataset
You need to re-do the measures
on a TEST dataset

```
-----trained linear model-----
Perf.           = 0.836647
Type of index = mse
MSE              = 0.836647
R2               = 0.976989
-----
```

Try a feedforward neural network!

```
net2 = feedforwardnet (3)
```

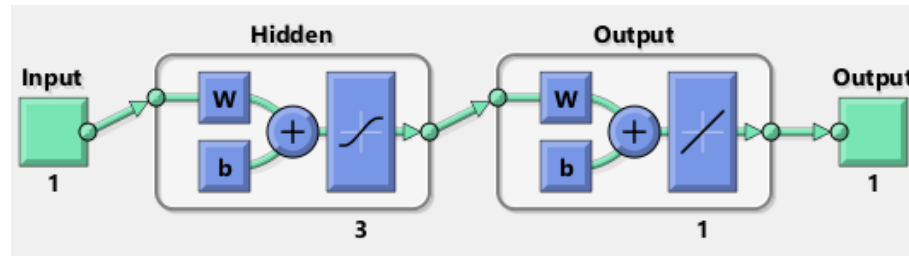
Initialization with 3 neurons

```
net2 = train(net2,x,t);
```

training

```
y2 = net2(x)
```

running



Try a feedforward neural network!

```
%% % Let's face the OCCAM's RAZOR --> use a large network
net2 = feedforwardnet(3);

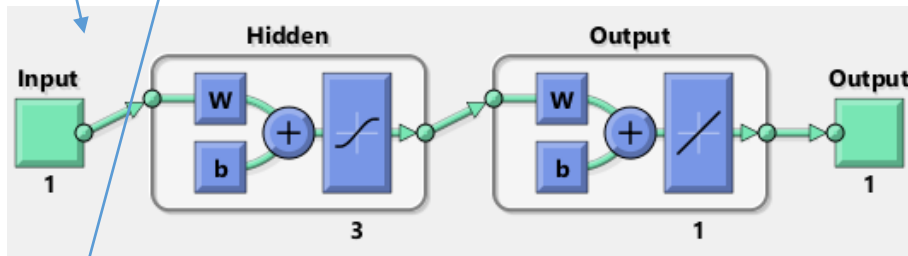
net2.divideParam.trainRatio = 1.0;
net2.divideParam.testRatio = 0.0;
net2.divideParam.valRatio = 0;

[net2 , tr] = train(net2,x,t);

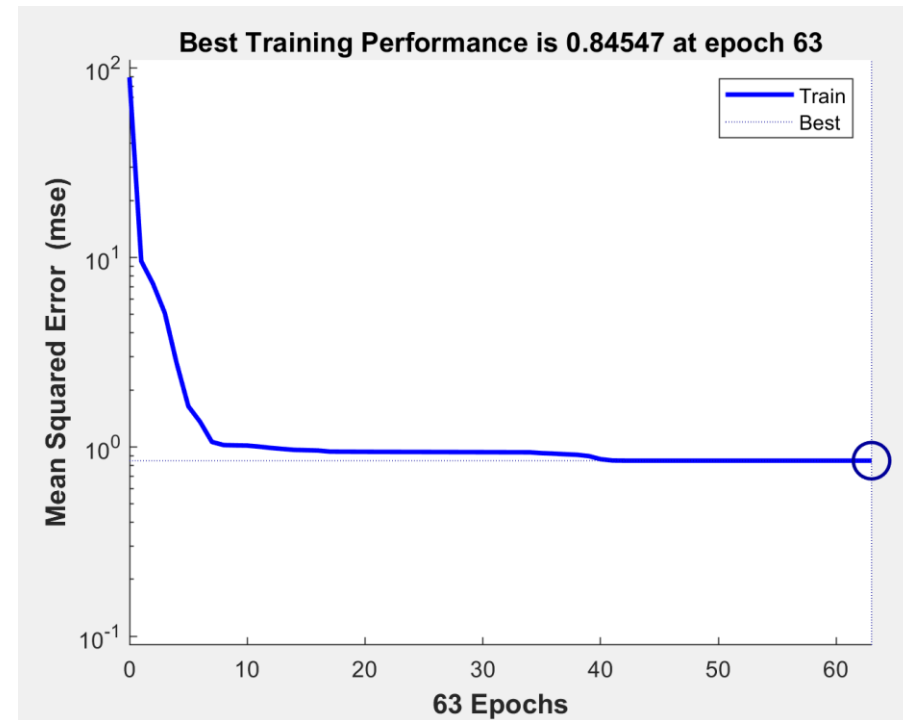
h2 = figure;
plotperform(tr)

view(net2)
y2 = net2(x);
perf = perform(net2,y2,t) % Error on **trainig**
```

Force to use
all data in
the TRAIN set



perf =
0.8455



MSE

Algorithms

Data Division: Random (dividerand)
Training: Levenberg-Marquardt (trainlm)
Performance: Mean Squared Error (mse)
Calculations: MEX

Progress

Epoch:	0	63 iterations	1000
Time:	0:00:00		
Performance:	89.2	0.845	0.00
Gradient:	179	8.16e-08	1.00e-07
Mu:	0.00100	1.00e-20	1.00e+10

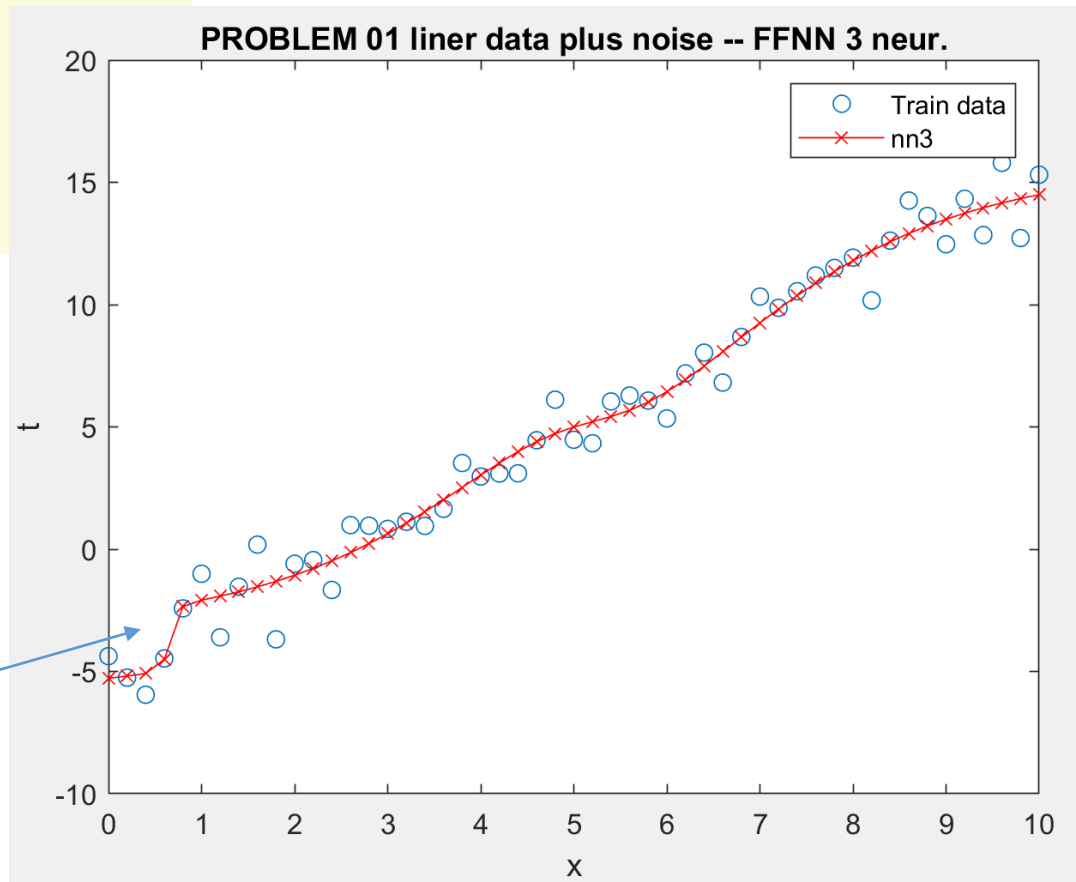
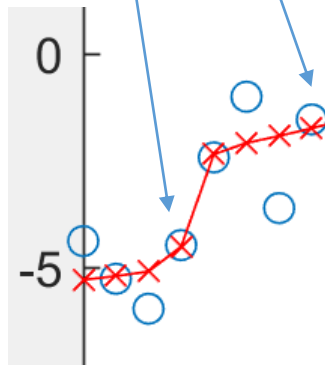
Try a feedforward neural network!

```
h3 = figure
plot(x,t, 'o');
title('PROBLEM 01 liner data plus noise -- FFNN 3 neur. ');
xlabel('x');
ylabel('t');

hold on
plot(x,y2, 'xr-');
xlabel('x');
ylabel('t');
legend('Train data', 'nn3')
```

Noise
fitting?

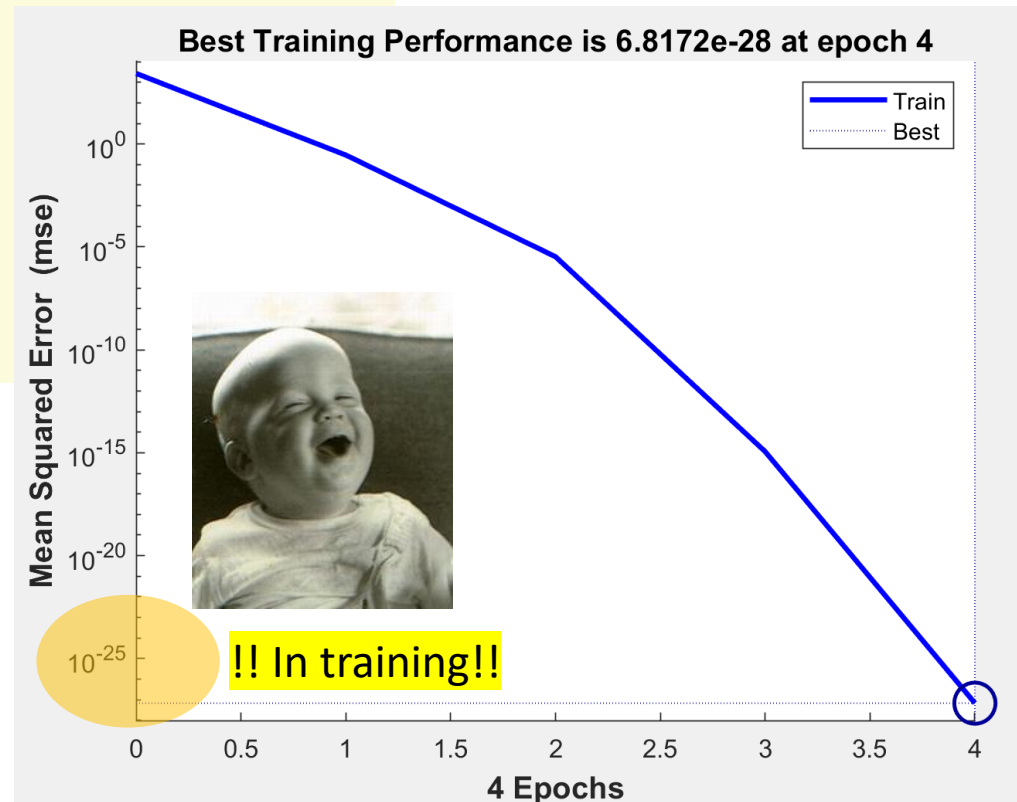
correct



Wrong design: 100 neurons with 51 data points!

```
%% real overfitting
% Let's face the OCCAM's RAZOR --> use very large network
net3 = feedforwardnet(100);
net3.divideParam.trainRatio = 1.0;
net3.divideParam.testRatio  = 0.0;
net3.divideParam.valRatio   = 0;

[net3 , tr] = train(net3,x,t);
figure(h2); plotperform(tr)
view(net3)
y3 = net3(x);
```

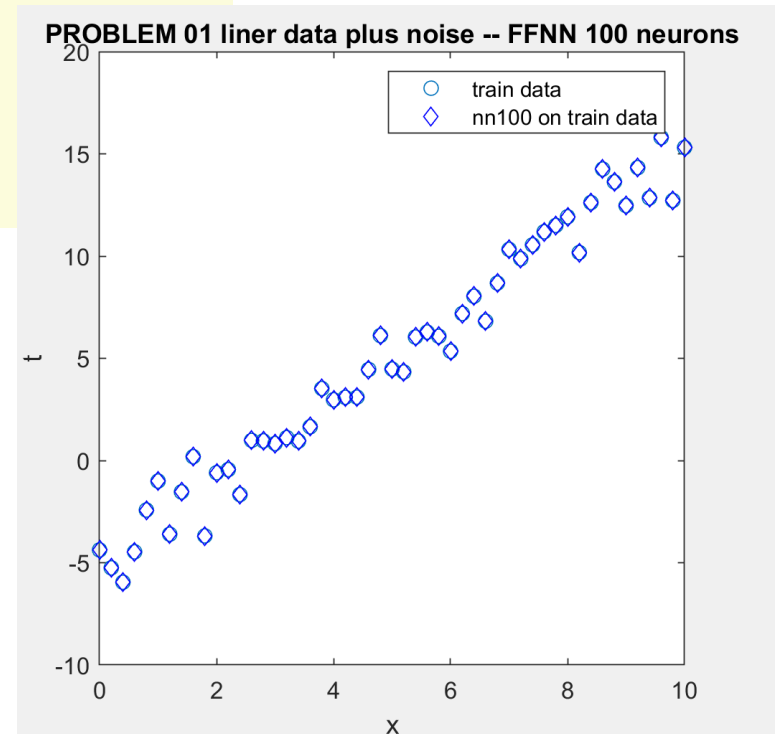
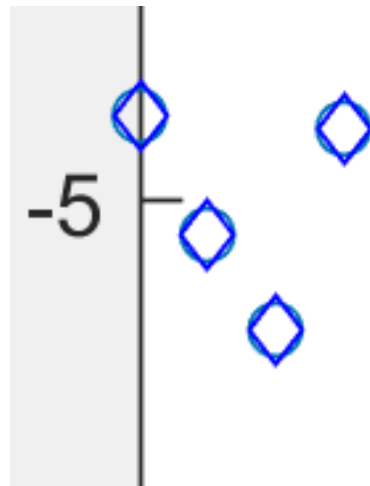


Wrong design: 100 neurons with 51 data points!

```
h4 = figure
subplot(1,2,1)
plot(x,t, 'o');
title('PROBLEM 01 liner data plus noise -- FFNN 100 neurons');
xlabel('x');
ylabel('t');
hold on
plot(x,y3, 'db');
legend('train data', 'nn100 on train data')
```



○ train data
◇ nn100 on train data



Let's built a test set....

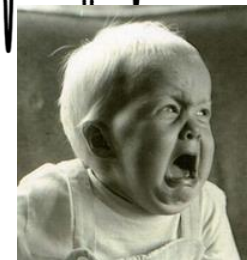
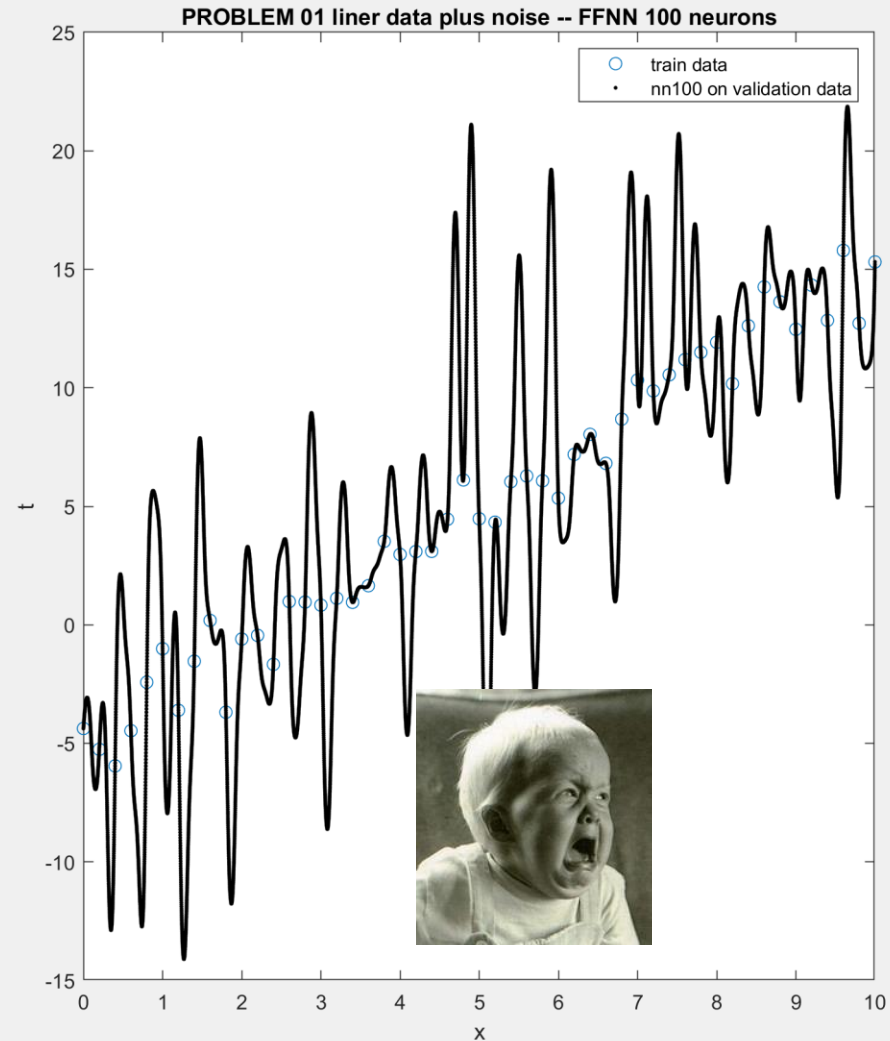
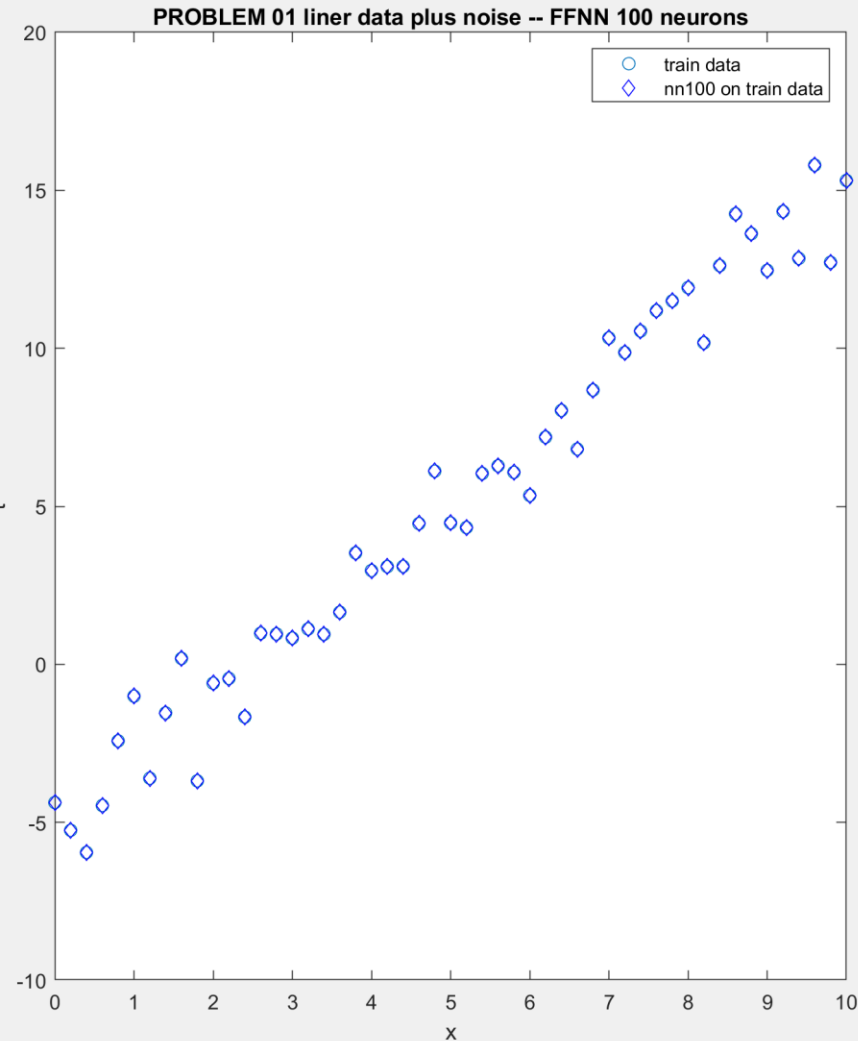
Let's inject in the net3 (100 neurons) more x points...

```
%% test point
x_bis = [0:0.0005:10];
y3_bis = net3(x_bis);

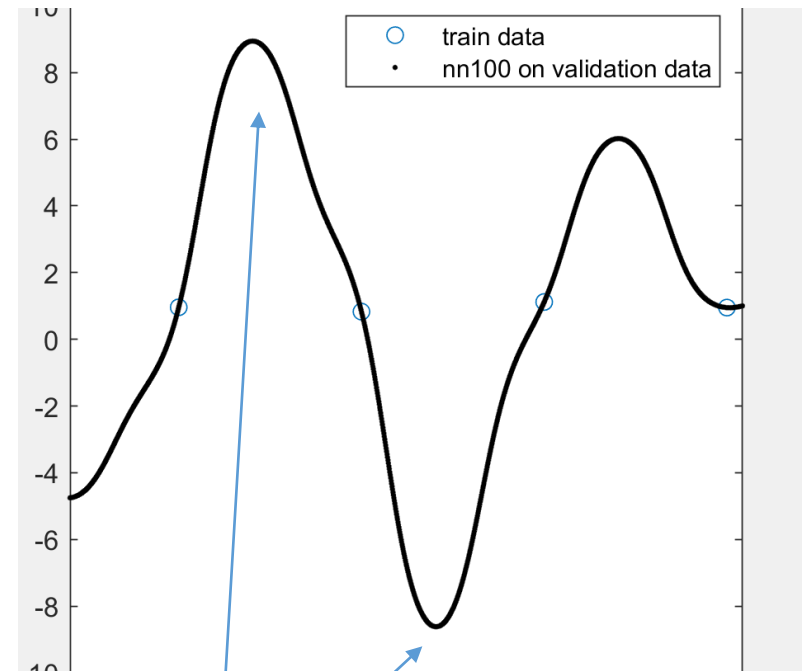
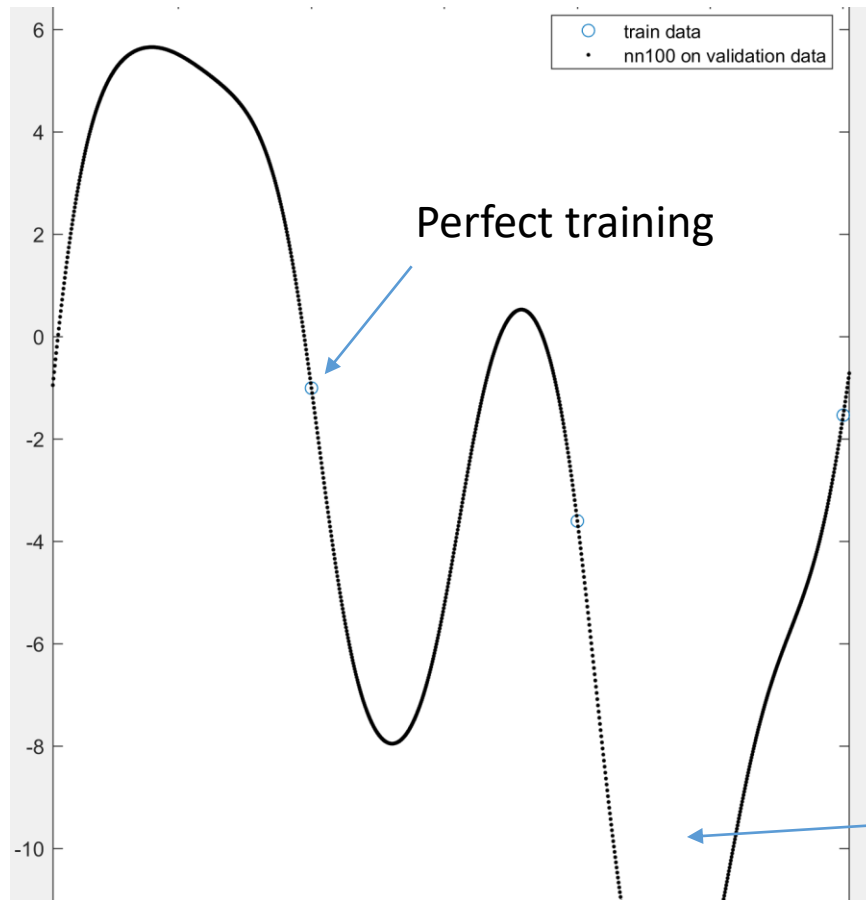
perf = perform(net3,y3,t) % Error on **trainig**

subplot(1,2,2)
plot(x,t, 'o');
title('PROBLEM 01 liner data plus noise -- FFNN 100 neurons');
xlabel('x');
ylabel('t');
hold on
plot(x_bis,y3_bis, '.k');
legend('train data', 'nn100 on validation data' )
```

...and we overfitted (a lot!)



Remember next time you will initialize a model with **#DoF > #samples...**



Very bad
generalization



Python file in Colab

- *1) Follow next slides*
- *2) and then lunch your script.*

`laboratory_COLAB_linear_regressors.ipynb`

Let's review how to
create a regressor,
check accuracy and
doing some plots

Import the needed libraries:

- plotting
- numerical processing

```
import matplotlib.pyplot as plt  
import numpy as np
```



It adds support for large, multi-dimensional arrays and matrices, along with a large collection of high-level mathematical functions to operate on these arrays.

Data preparation

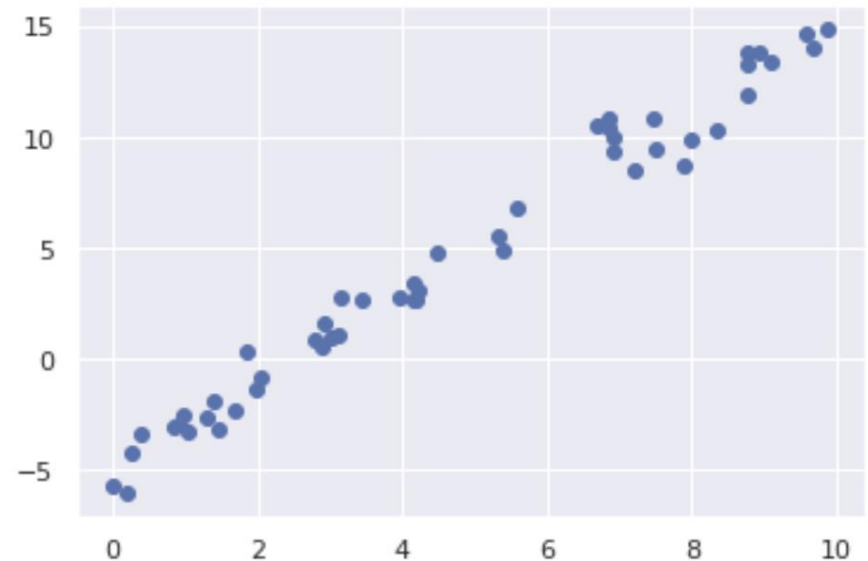
1D Linear Regression

We want to create a model which is a straight-line fit of the form

$$y = ax + b$$

Consider the following data, which is scattered about a line with a slope of 2 and an intercept of -5:

```
[35] rng = np.random.RandomState(1)
      num_points = 50
      x = 10 * rng.rand(num_points)
      t = 2 * x - 5 + rng.randn(num_points)
      plt.scatter(x, t);
```



Creation of the linear regressor

We can use **Scikit-Learn**'s LinearRegression estimator to fit this data and construct the best-fit line:

```
from sklearn.linear_model import LinearRegression
model = LinearRegression(fit_intercept=True)

model.fit(x[:, np.newaxis], t) # newaxis adds a dimension (needed from the function)

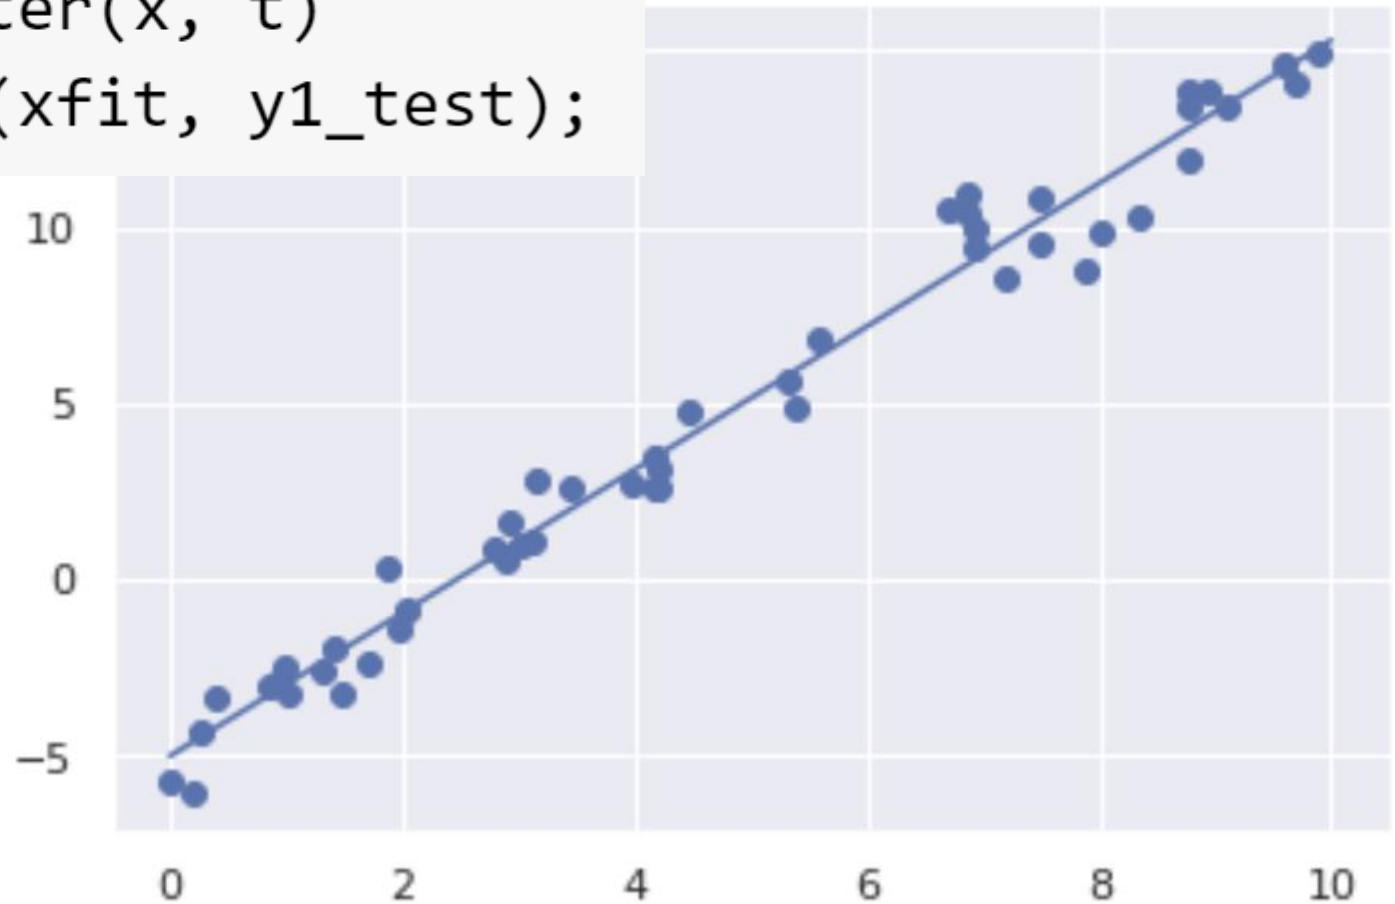
y1 = model.predict(x[:, np.newaxis]) # let's create the predicted values

xfit = np.linspace(0, 10, 1000) # let's create an array to test the regressor [0 ... 10] of 1000 points
y1_test = model.predict(xfit[:, np.newaxis]) # let's create the predicted values
```

```
model = LinearRegression(fit_intercept=True)
model.fit(x[:, np.newaxis], t)
y1 = model.predict(x[:, np.newaxis])
```

Plotting

```
plt.scatter(x, t)  
plt.plot(xfit, y1_test);
```



Get the parameters of the model

```
[37] print("Model slope:      ", model.coef_[0])  
      print("Model intercept:", model.intercept_)
```

```
↳ Model slope:      2.0272088103606953  
   Model intercept: -4.998577085553204
```

Accuracy assessement

```
[38] from sklearn.metrics import mean_squared_error
      from math import sqrt

      # Mean squared error
      mse = mean_squared_error(y1, t)
      rmse = sqrt(mse)

      #R^2
      from sklearn.metrics import r2_score
      R2 = r2_score(y1, t)

      print("MSE:      ", mse )
      print("RMSE:     ", rmse )
      print("R2:       ", R2 )
```

```
☞ MSE:      0.8183388570266171
   RMSE:     0.9046208360559783
   R2:       0.9786330659856474
```

**** WARNING ****
You have to create a TEST
dataset to get the real
accuracy in generalization.
This is processed on the
TRAIN dataset.

You will find different numbers w.r.t. Matlab
due to the different data initialization

Main points



- *New model? What to know? What to do?*
- *Creation and use (with code) of the first complete machine learning model*
 - *Linear regressors*
 - *First neural network*
 - *Accuracy assessment*

```
net1 = newlind(x,t)
```

```
y1 = net1(x)
```

```
net2 = feedforwardnet (3)
```

```
net2 = train(net2,x,t);
```

```
y2 = net2(x)
```

